

Confidential Computing with OpenBSD — The Next Step

Hans-Jörg Höxer

Confidential Computing with OpenBSD

Agenda

- **Introduction**
- First step: Memory encryption for VMs — SEV
- Next step: vCPU state encryption — SEV-ES
- Conclusion

About

Hans-Jörg Höxer

- Mid-2000s:
 - hshoexer@openbsd.org
- genua GmbH (www.genua.de):
 - hshoexer@genua.de
 - OpenBSD based products
 - Firewalls and VNP-Appliances
 - Confidential Computing

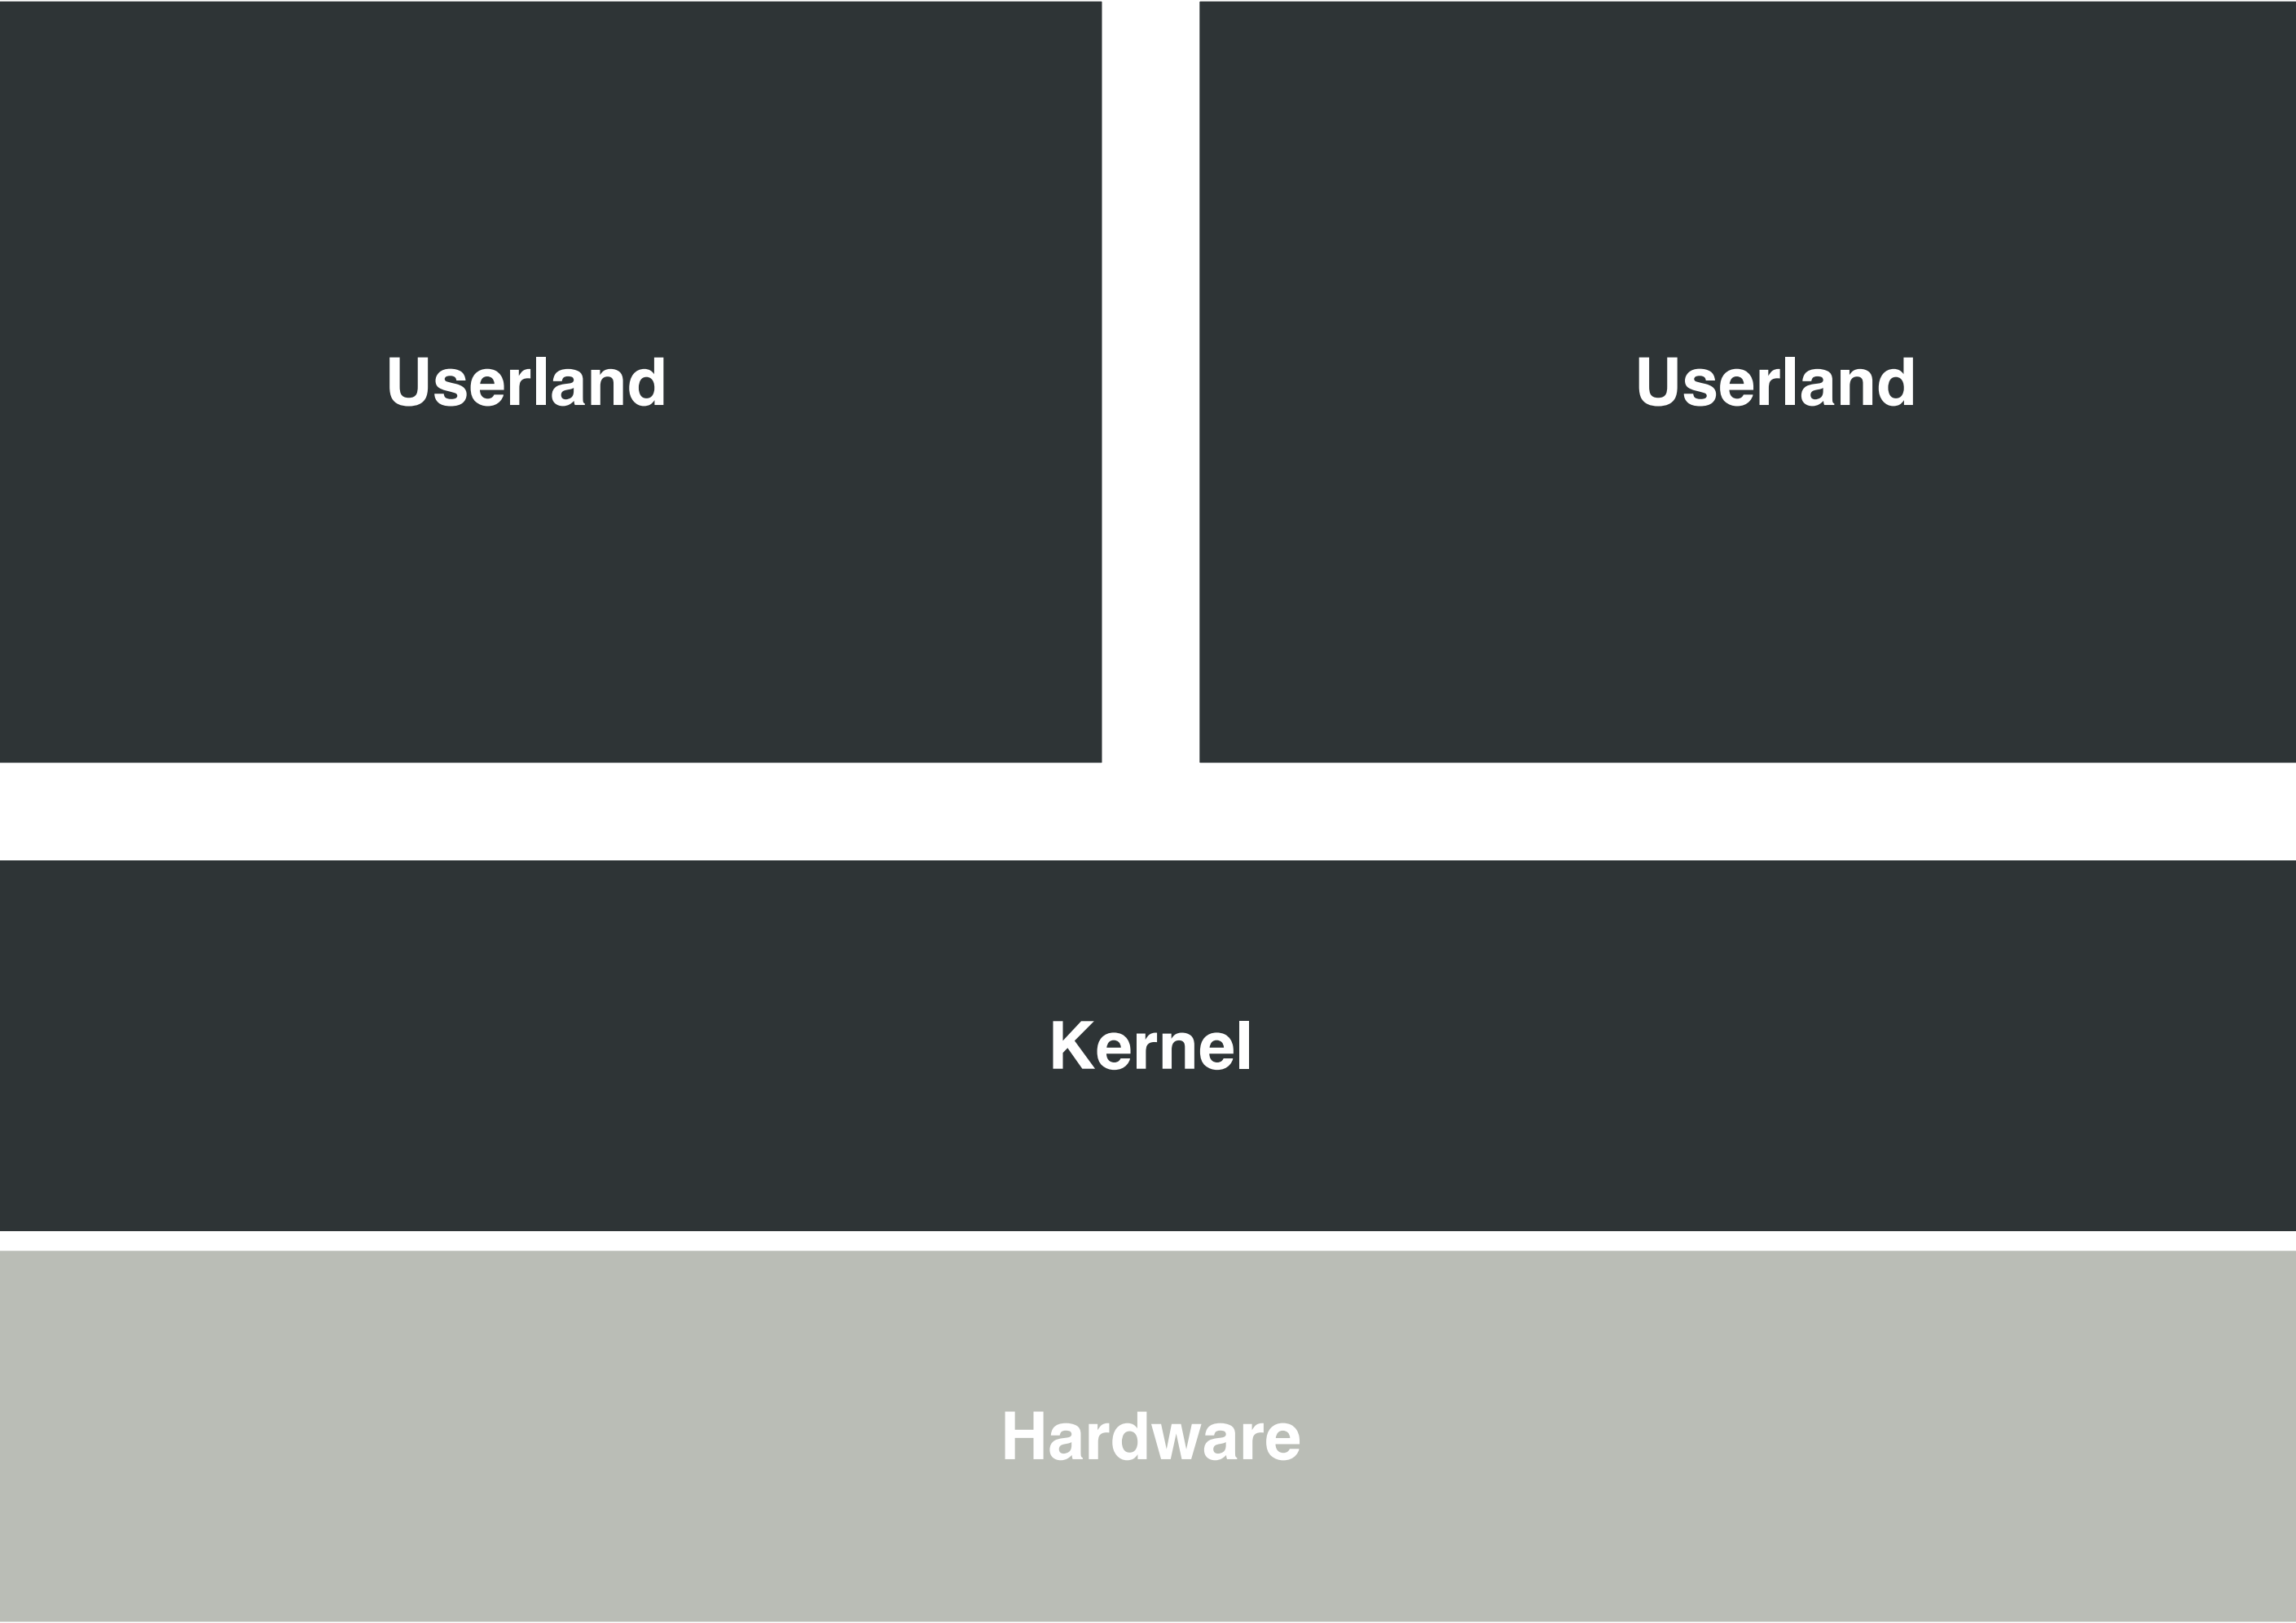
Confidential Computing

What is this all about?

- Problem:
 - Sensitive data in an untrusted environment
 - Context: Virtualisation, VMs, cloud
- Supposed solution:
 - “Turn public cloud into private cloud”
 - Bold claims...

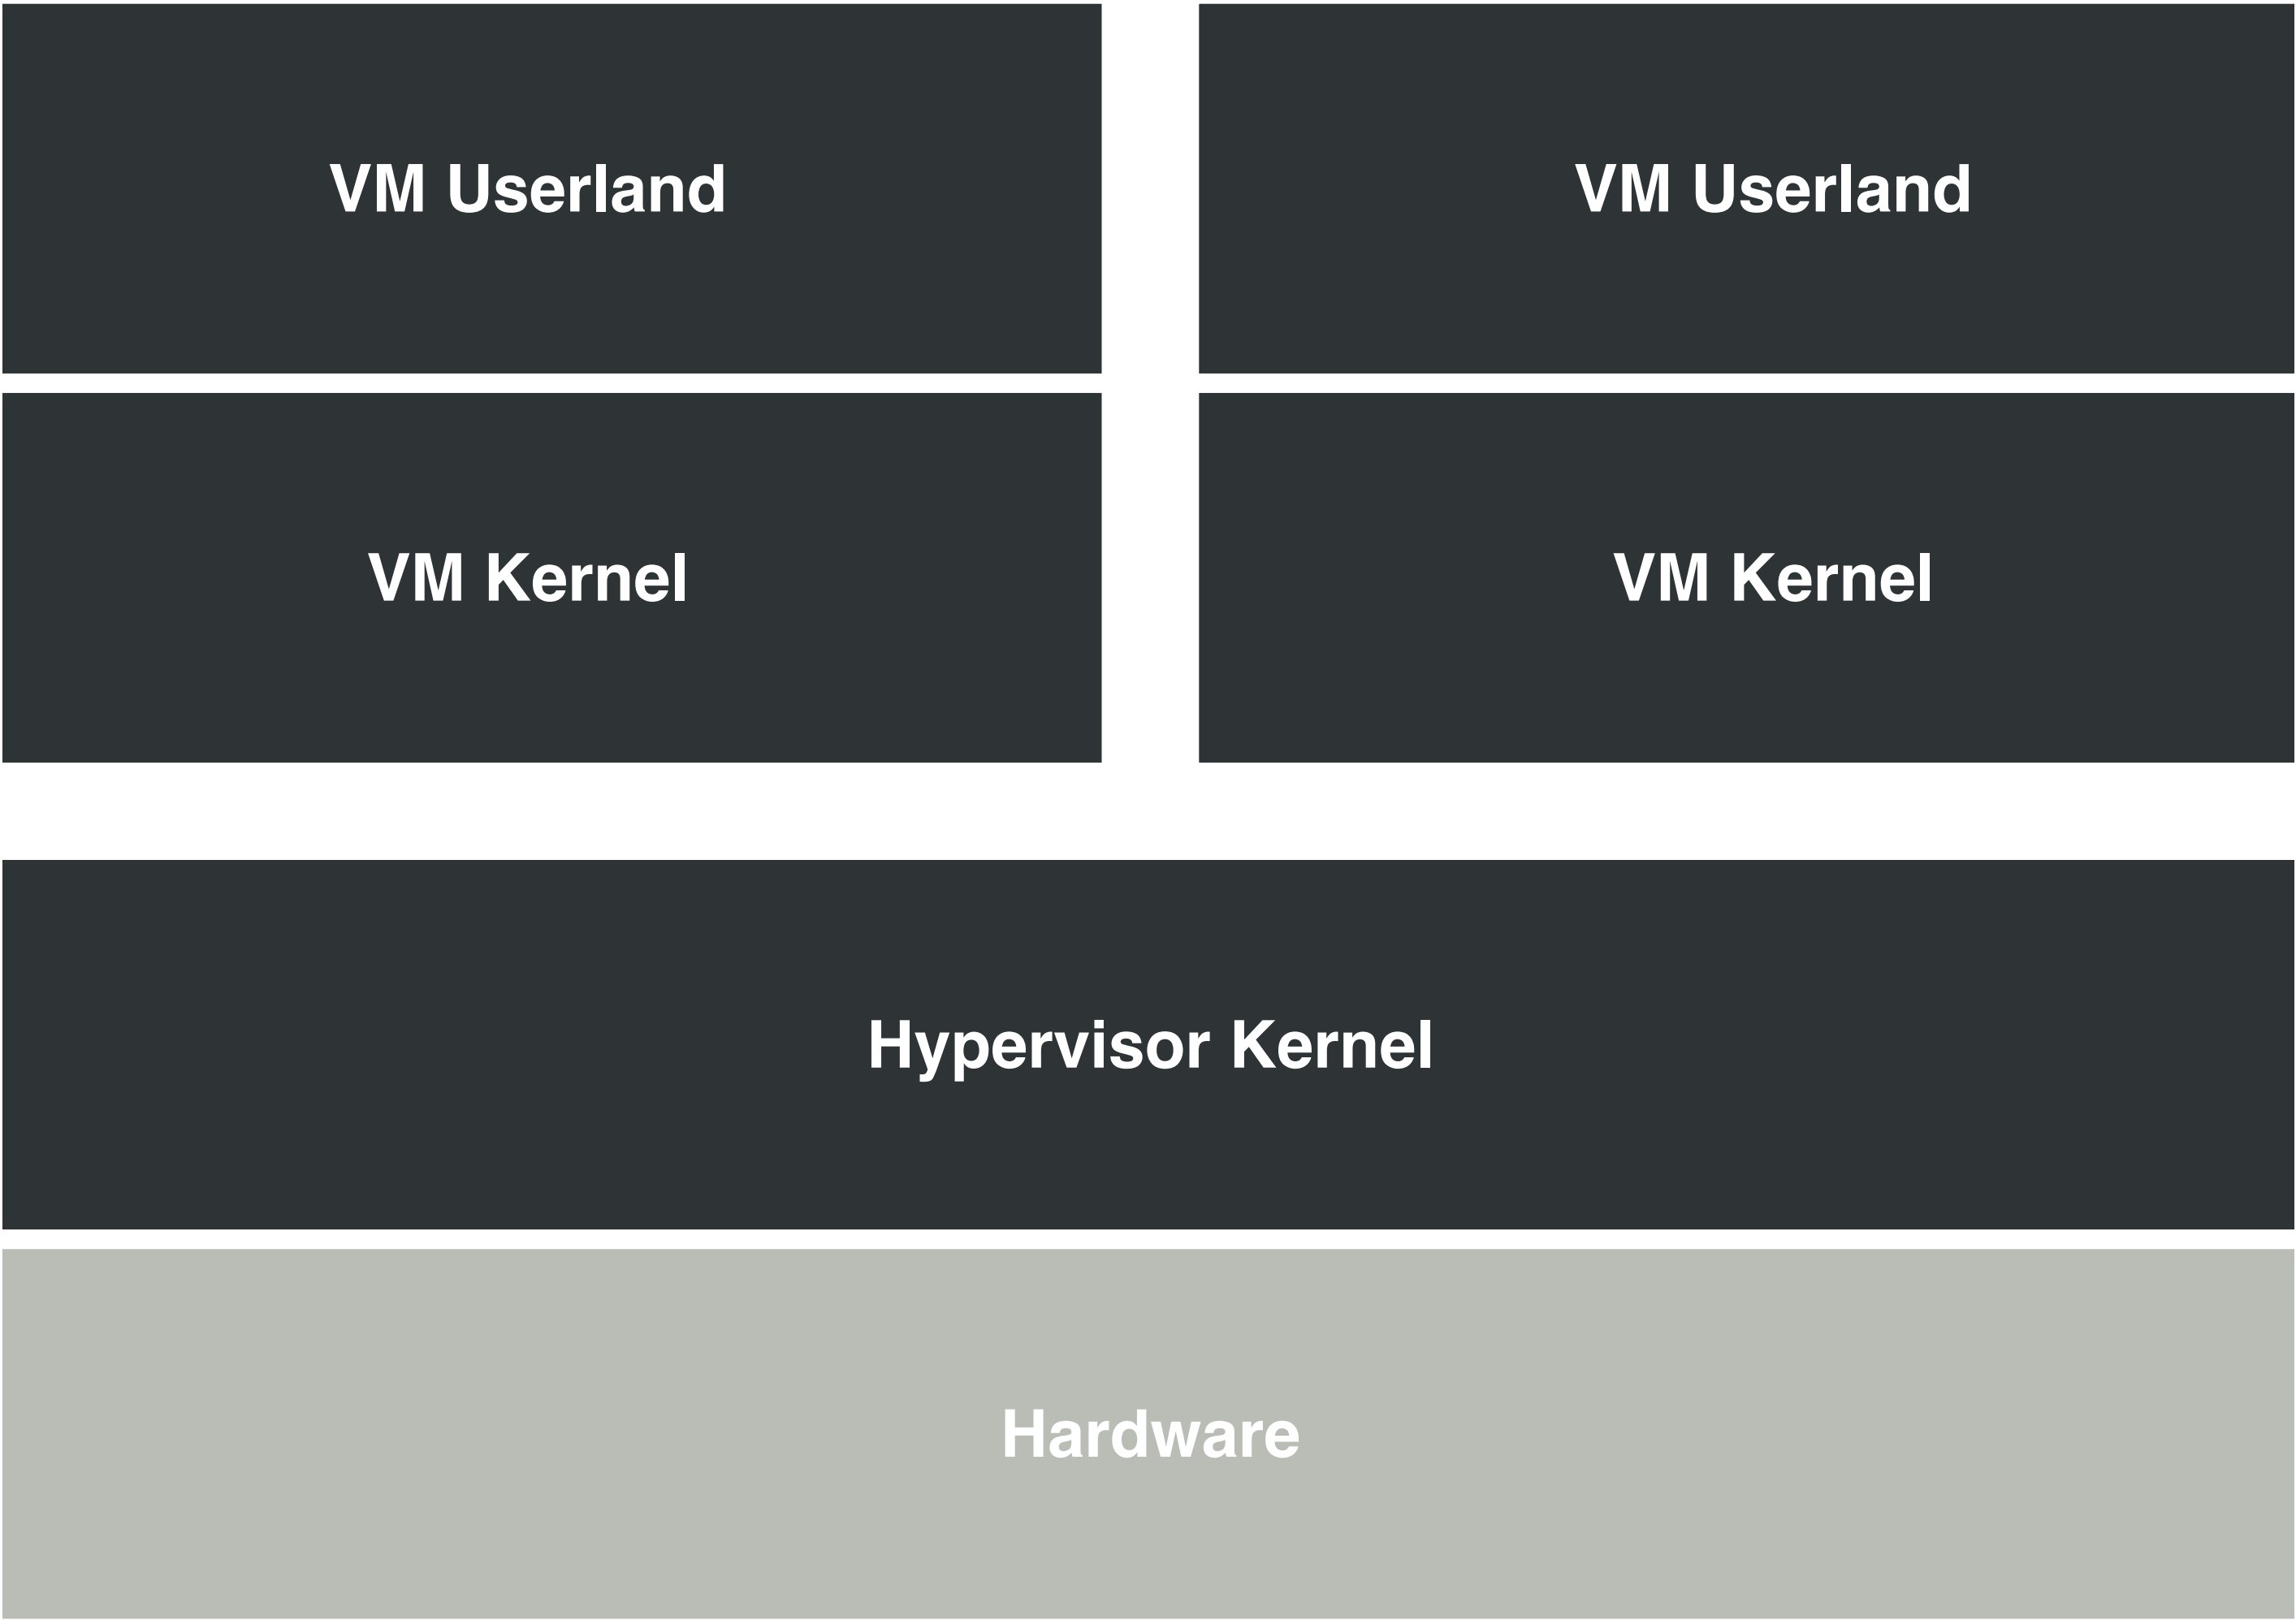
➡ Learn by implementing for OpenBSD

Untrusted Environments

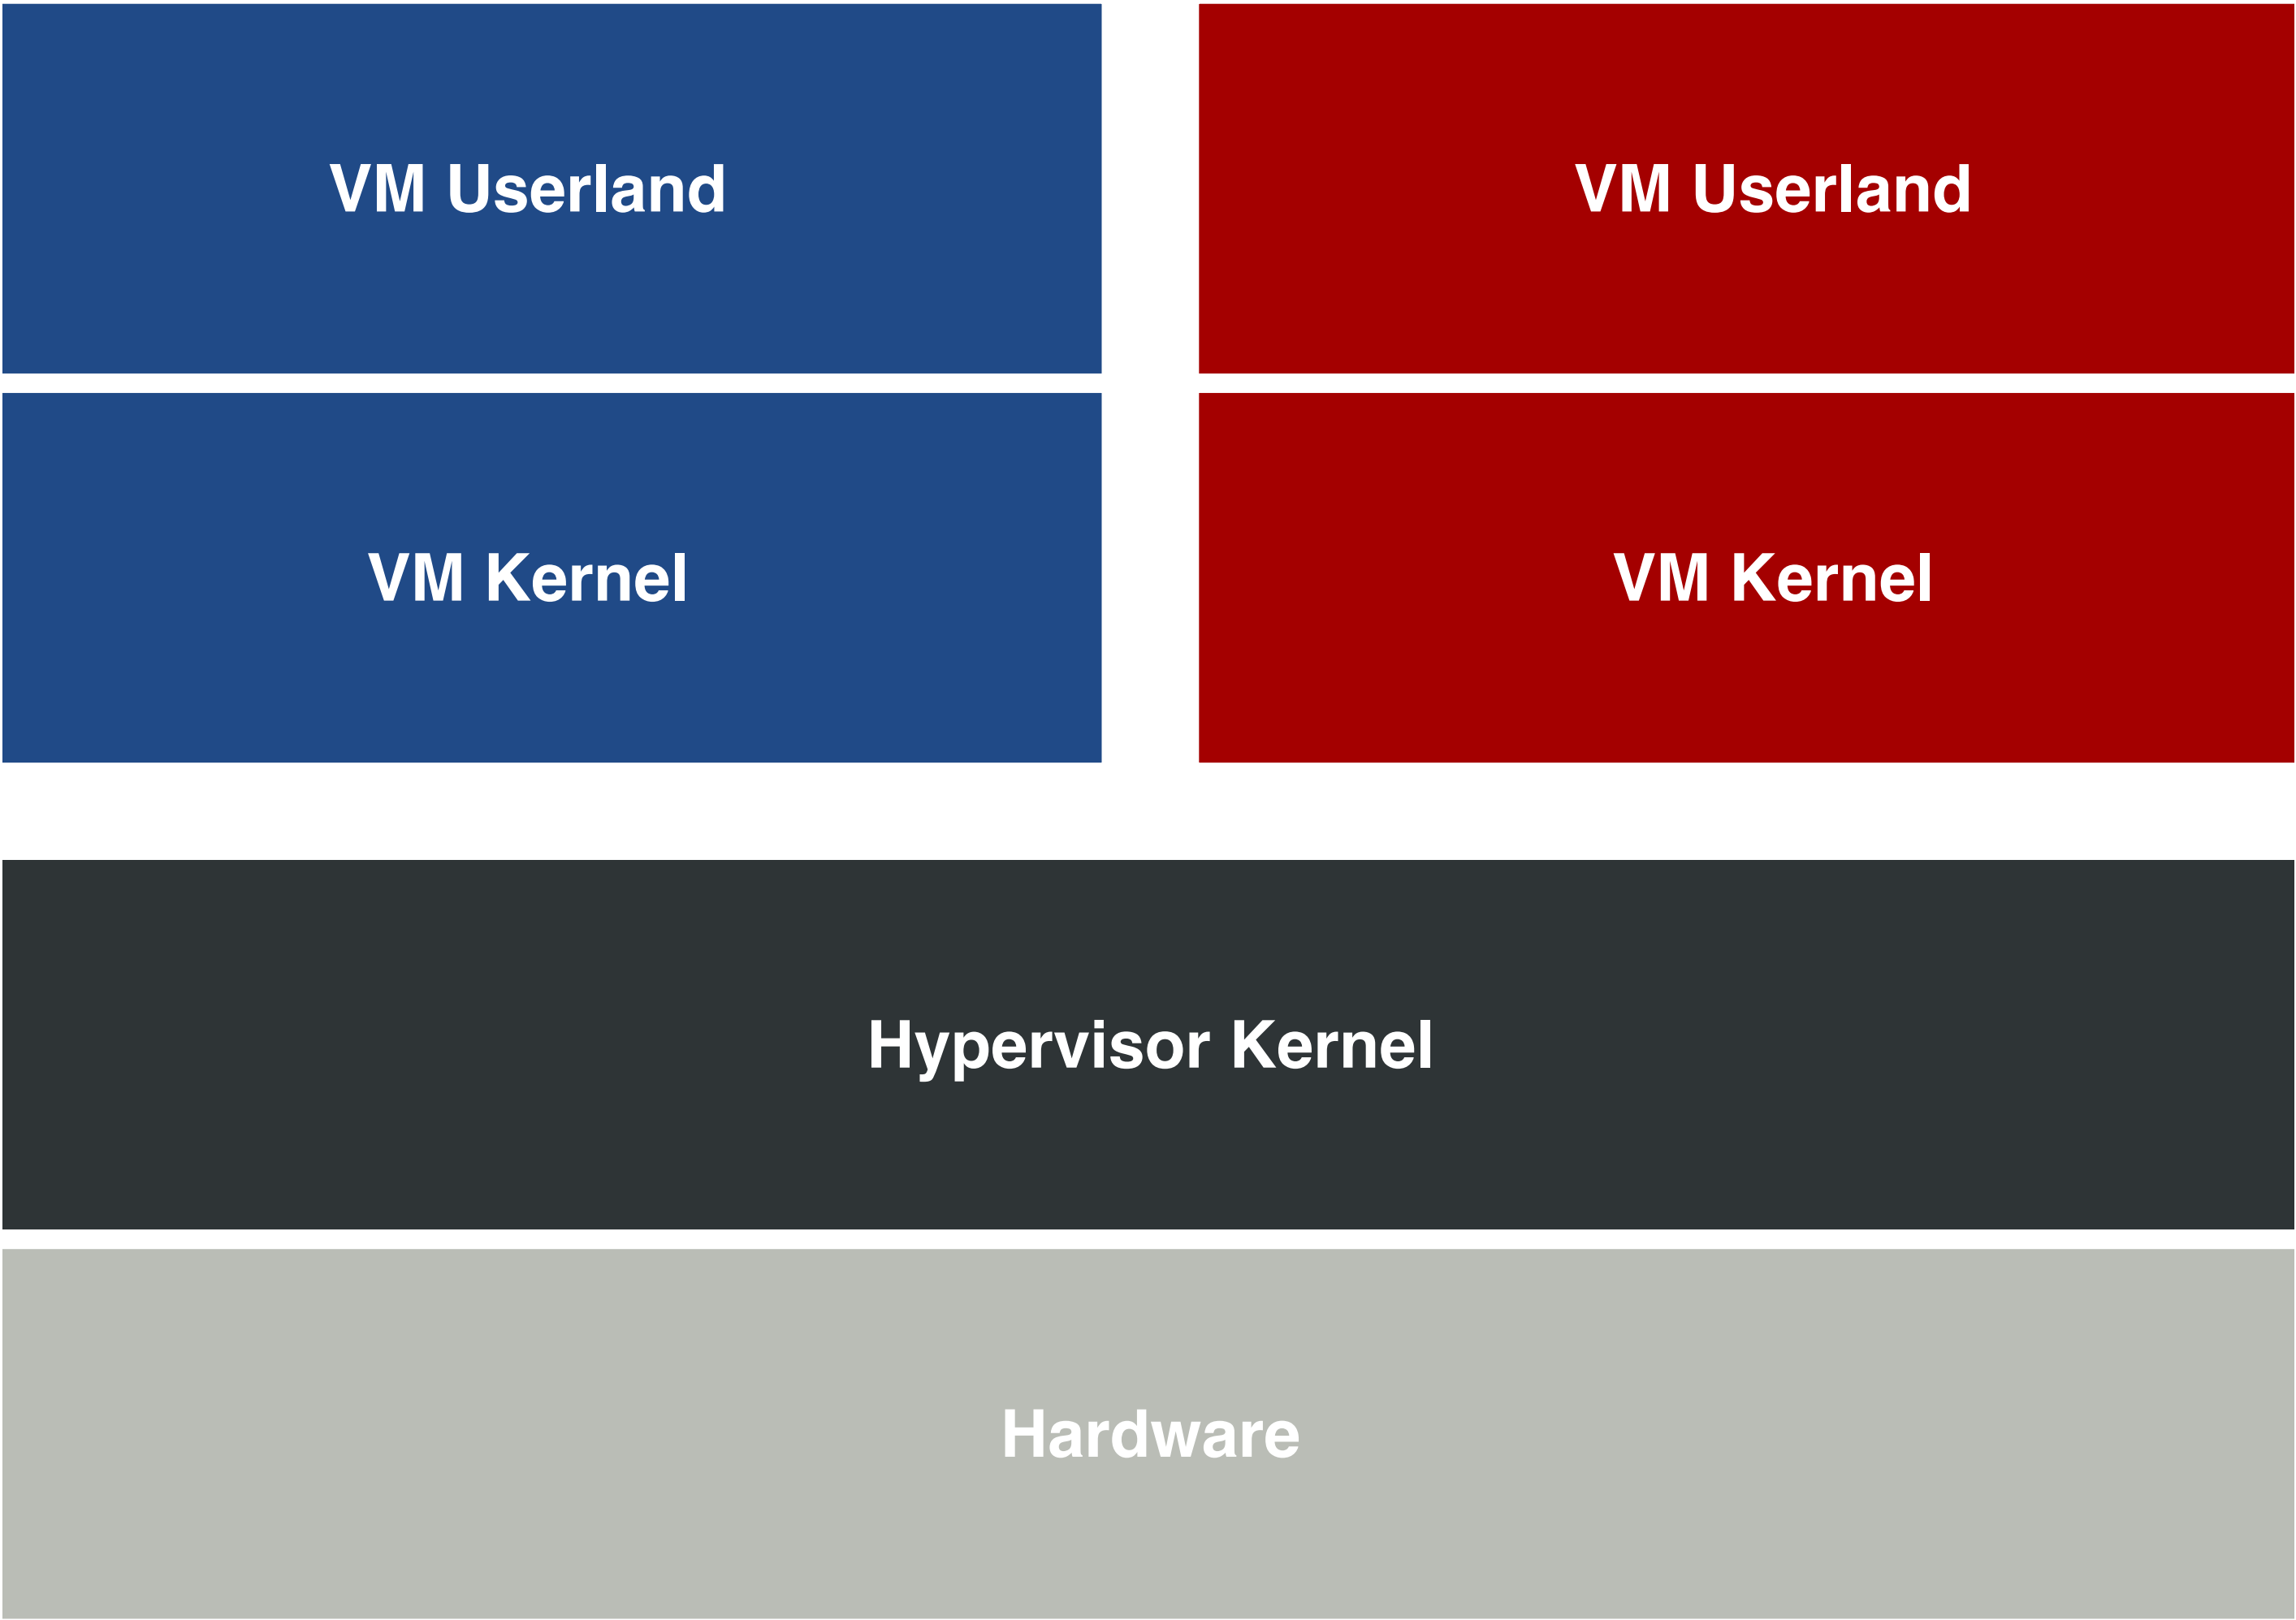


Generic OS

Untrusted Environments



Untrusted Environments



Confidential Computing

Claims

- Techniques to protect computing workload from its untrusted environment
 - Data confidentiality
 - Data integrity
 - Code integrity
- Isolation levels
 - Function or library isolation
 - Application isolation
 - ★ Virtual machine isolation

Confidential Computing with OpenBSD

Agenda

- Introduction
- **First step: Memory encryption for VMs — SEV**
- Next step: vCPU state encryption — SEV-ES
- Conclusion

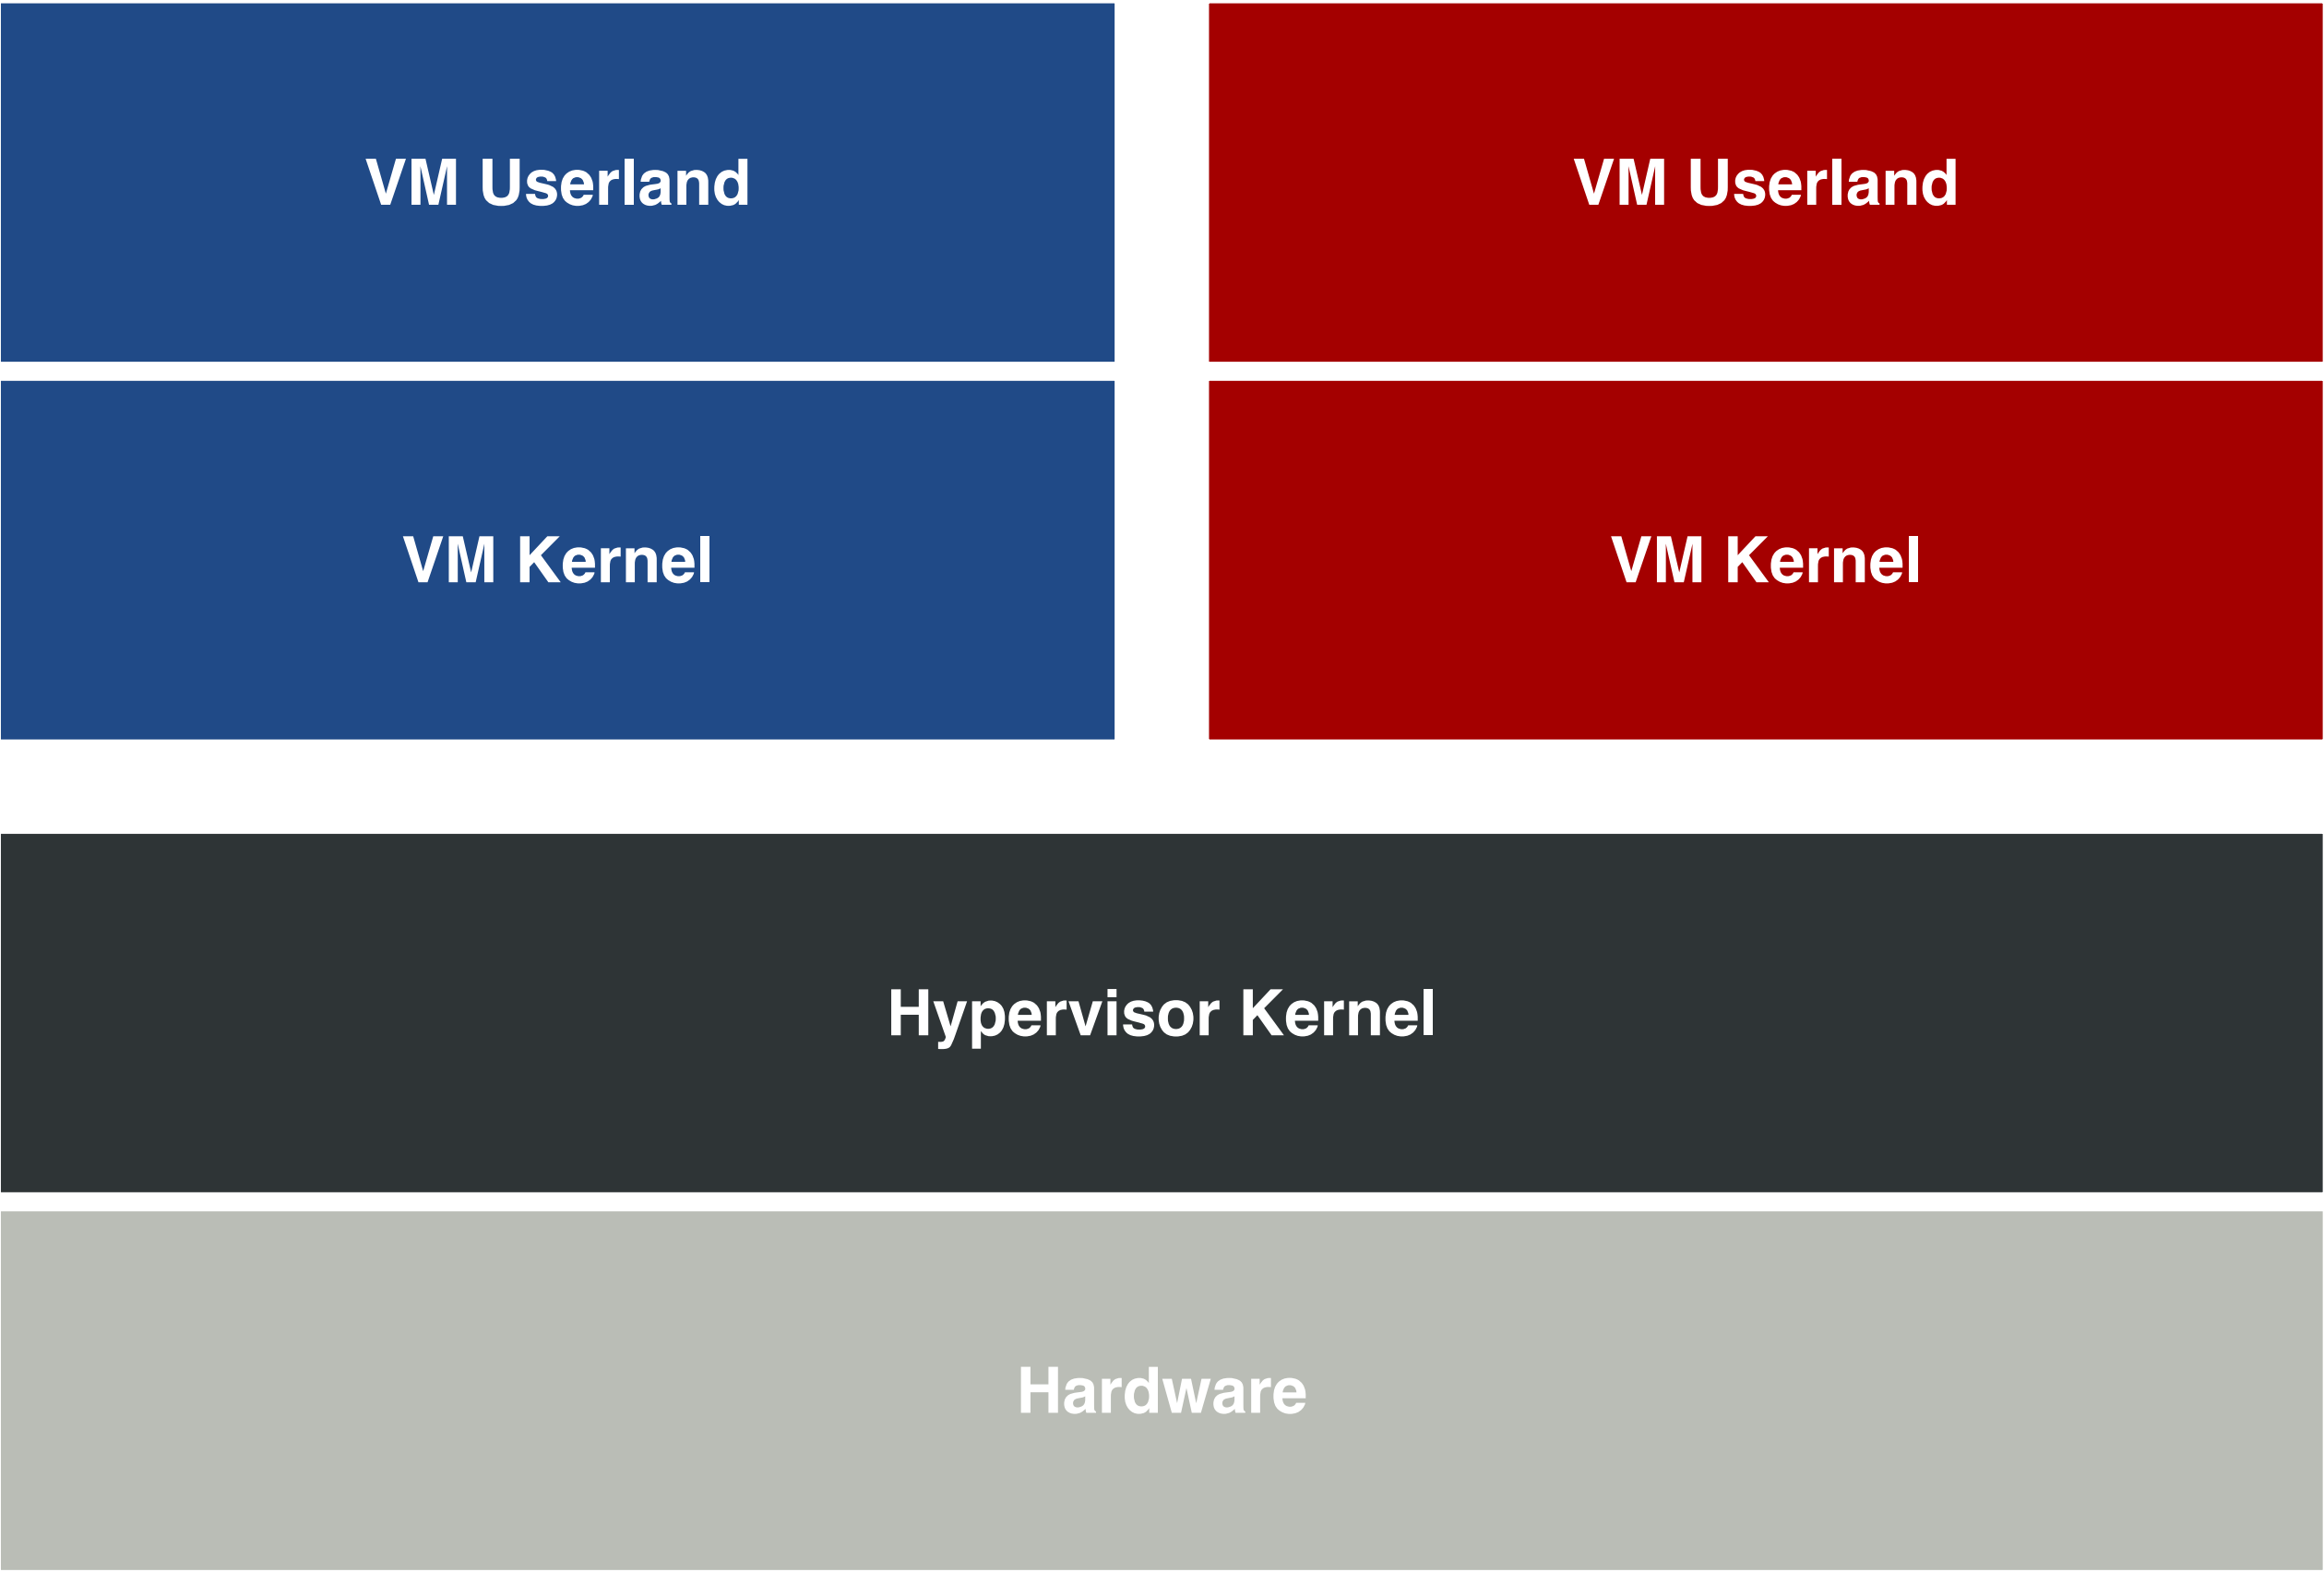
Confidential Computing

Hardware Support

- Hardware support:
 - ★ Runtime encryption
 - Attestation
 - Strong isolation
- Examples:
 - **AMD SEV, SEV-ES, SEV-SNP**
 - Intel TDX, Arm CCA

AMD Secure Encrypted Virtualisation

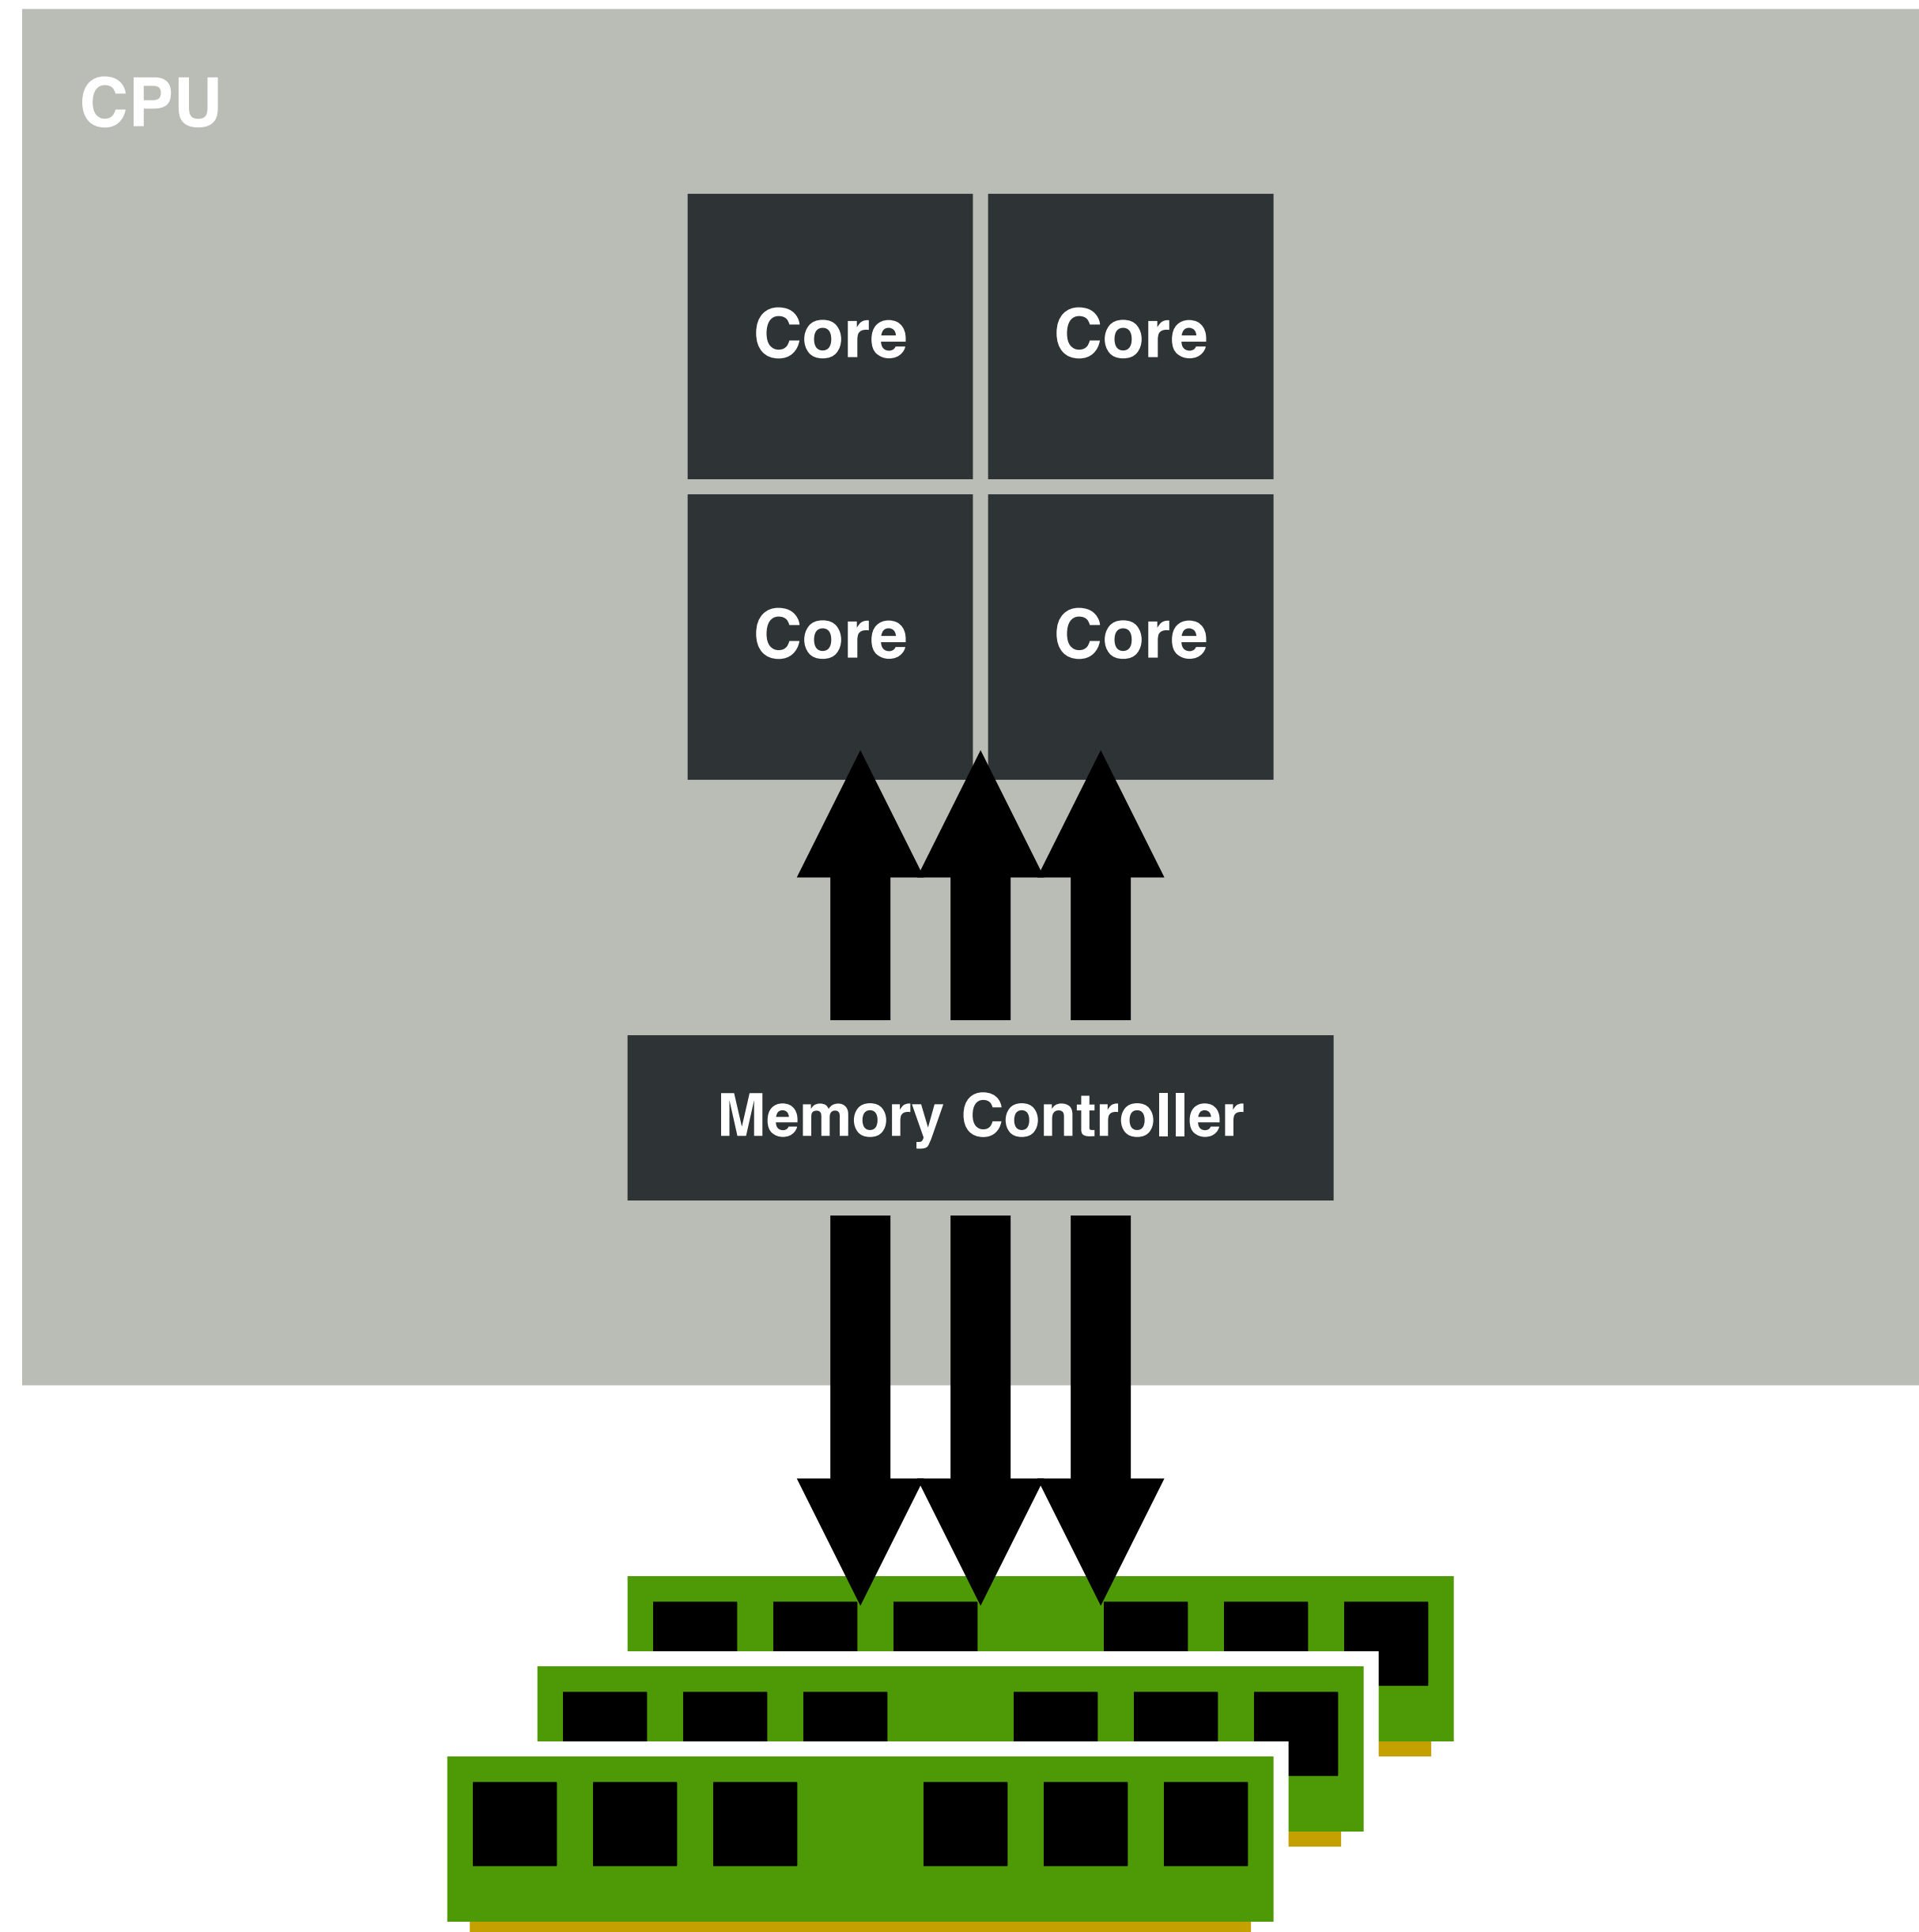
Confidential VM



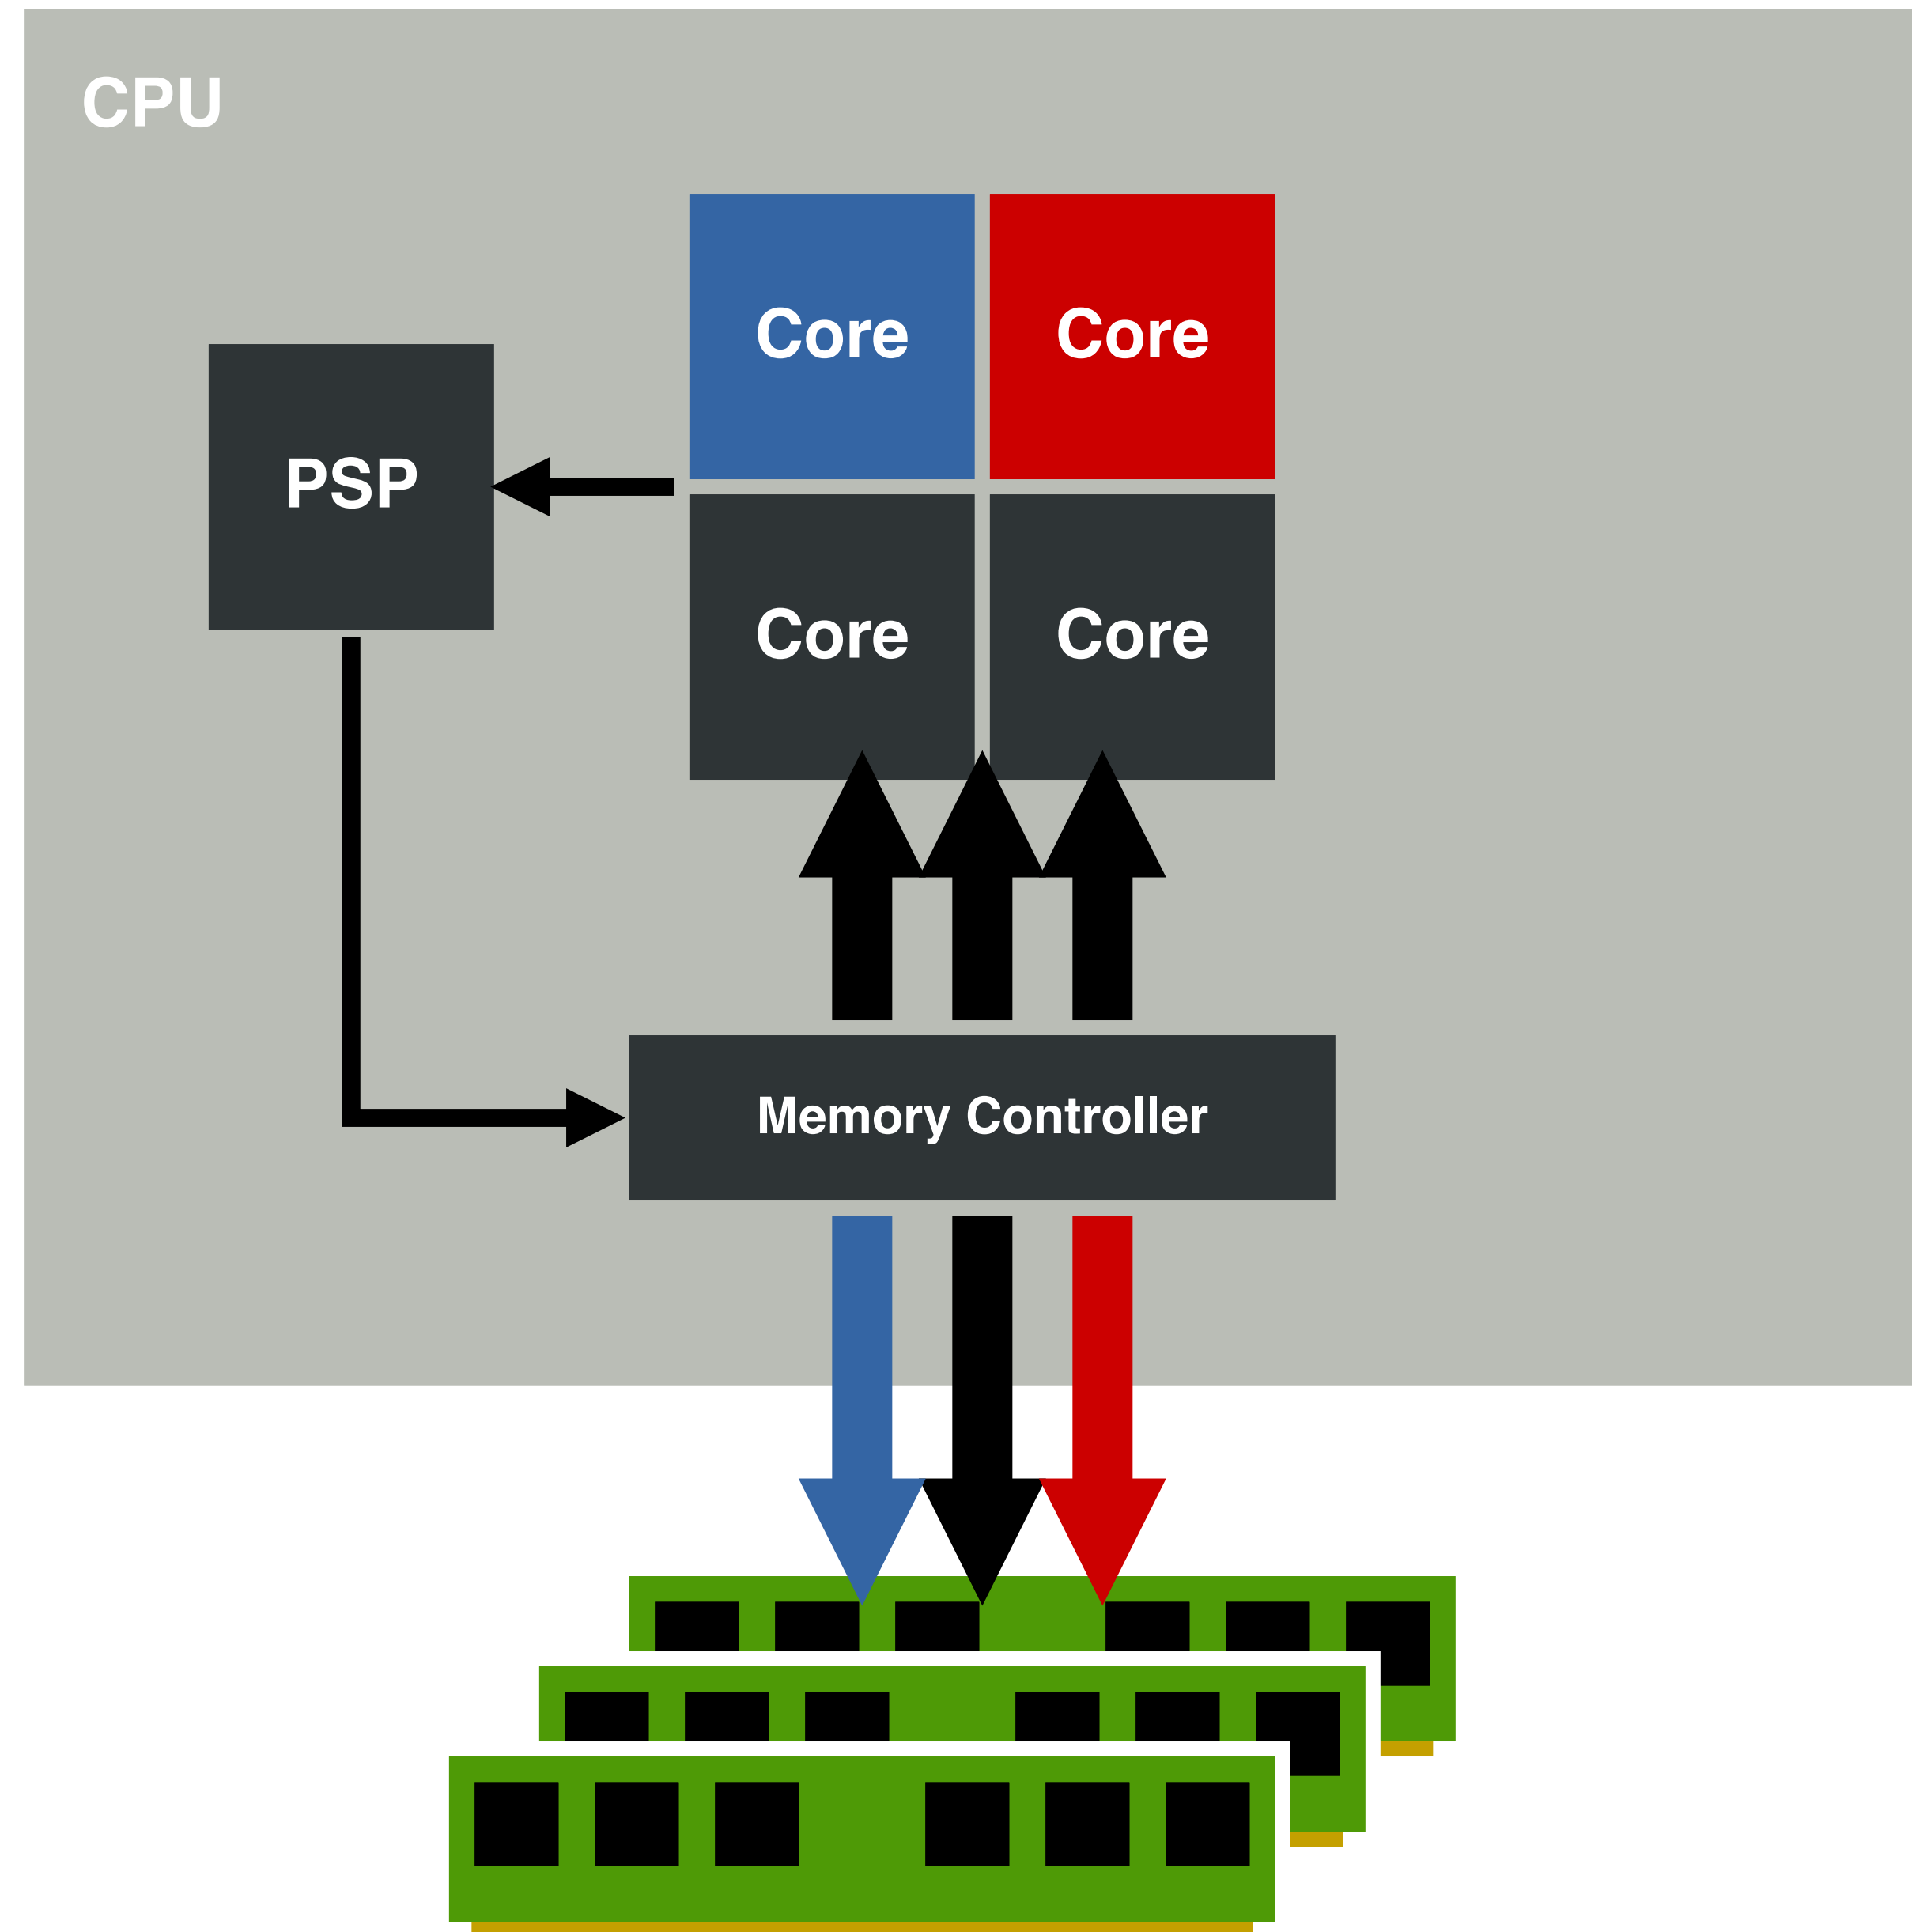
Confidential VM

AMD SEV

Architecture



AMD SEV Architecture



Confidential Computing for OpenBSD

Goals

- Implement support for AMD SEV-^{*}:
 - psp(4), vmd(8), vmm(4), GENERIC
 - Both host and guest
- Step by step:
 - ☒ SEV — OpenBSD 7.6 (October 2024)
 - ☐ SEV-ES — -current
 - ☐ SEV-SNP — work in progress
- Compatibility:
 - Linux/KVM host

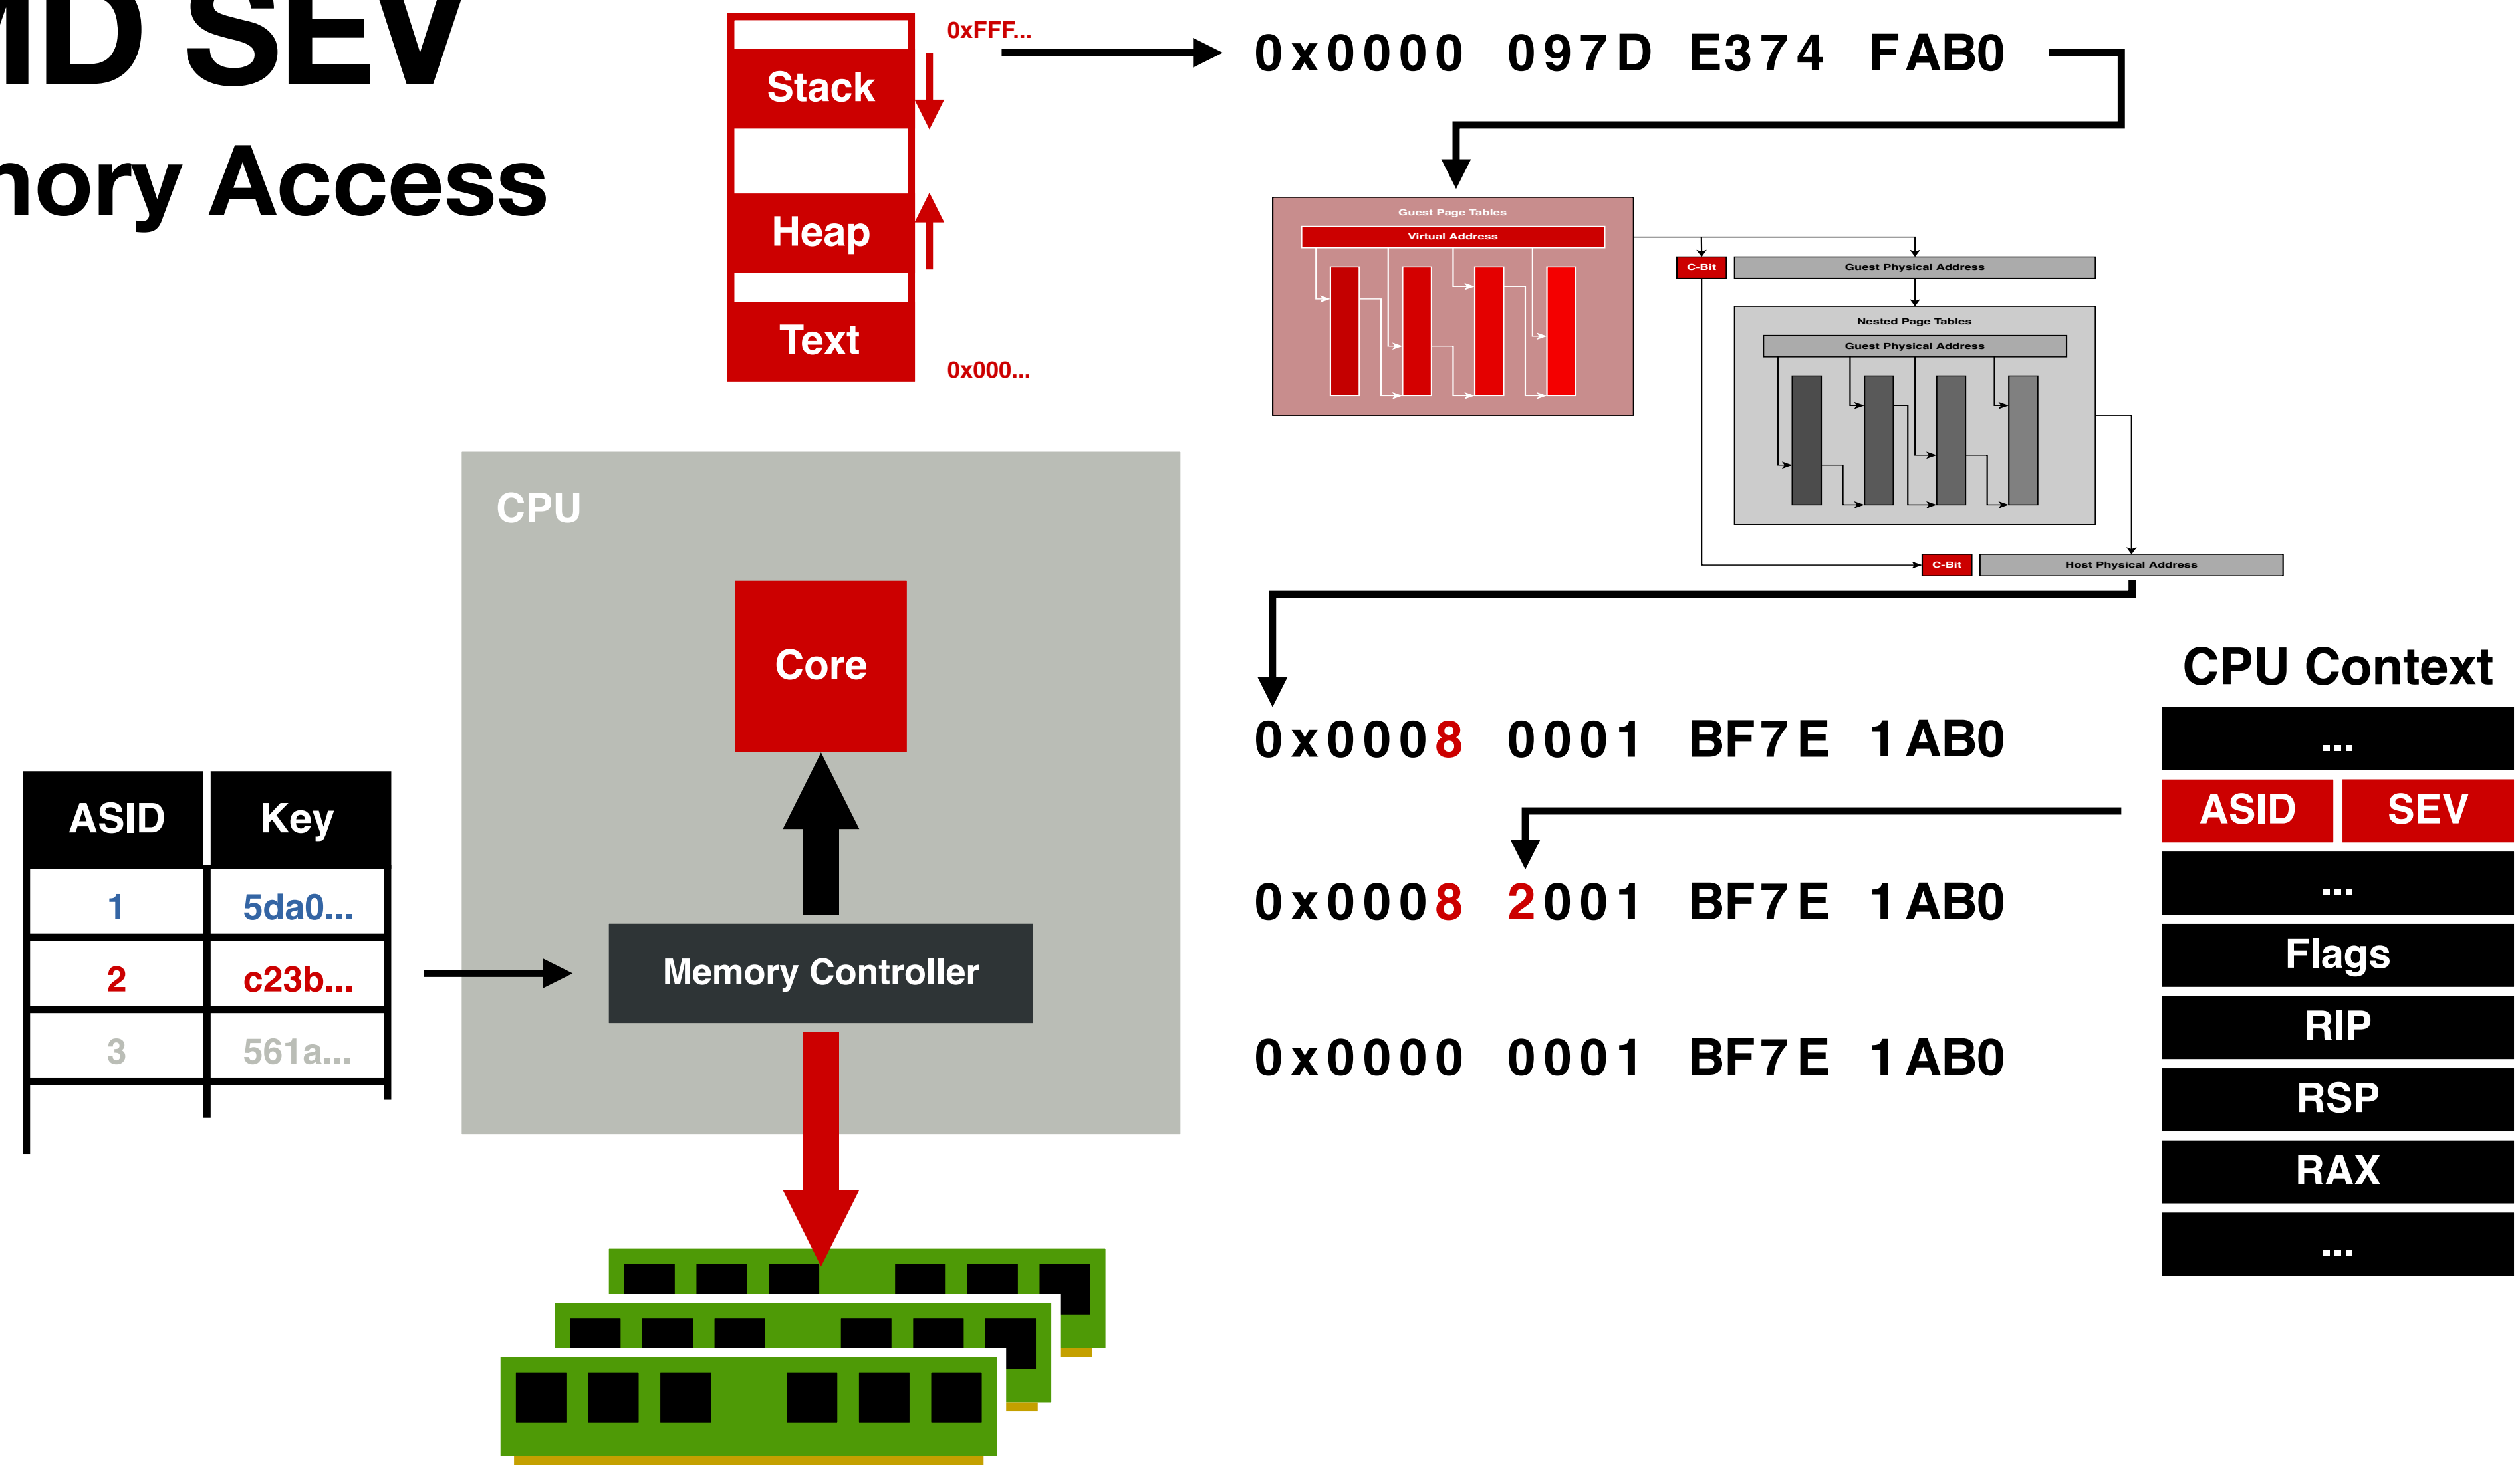
AMD SEV

Secure Encrypted Virtualisation

- Guest VM controls encryption!
 - Page tables:
 - “Crypt bit” (C-bit)
 - Private data
 - Public data — shareable:
 - DMA bounce buffers used by virtio(4)
 - Implemented in bus_dma(9)
- Guest and host support in OpenBSD 7.6

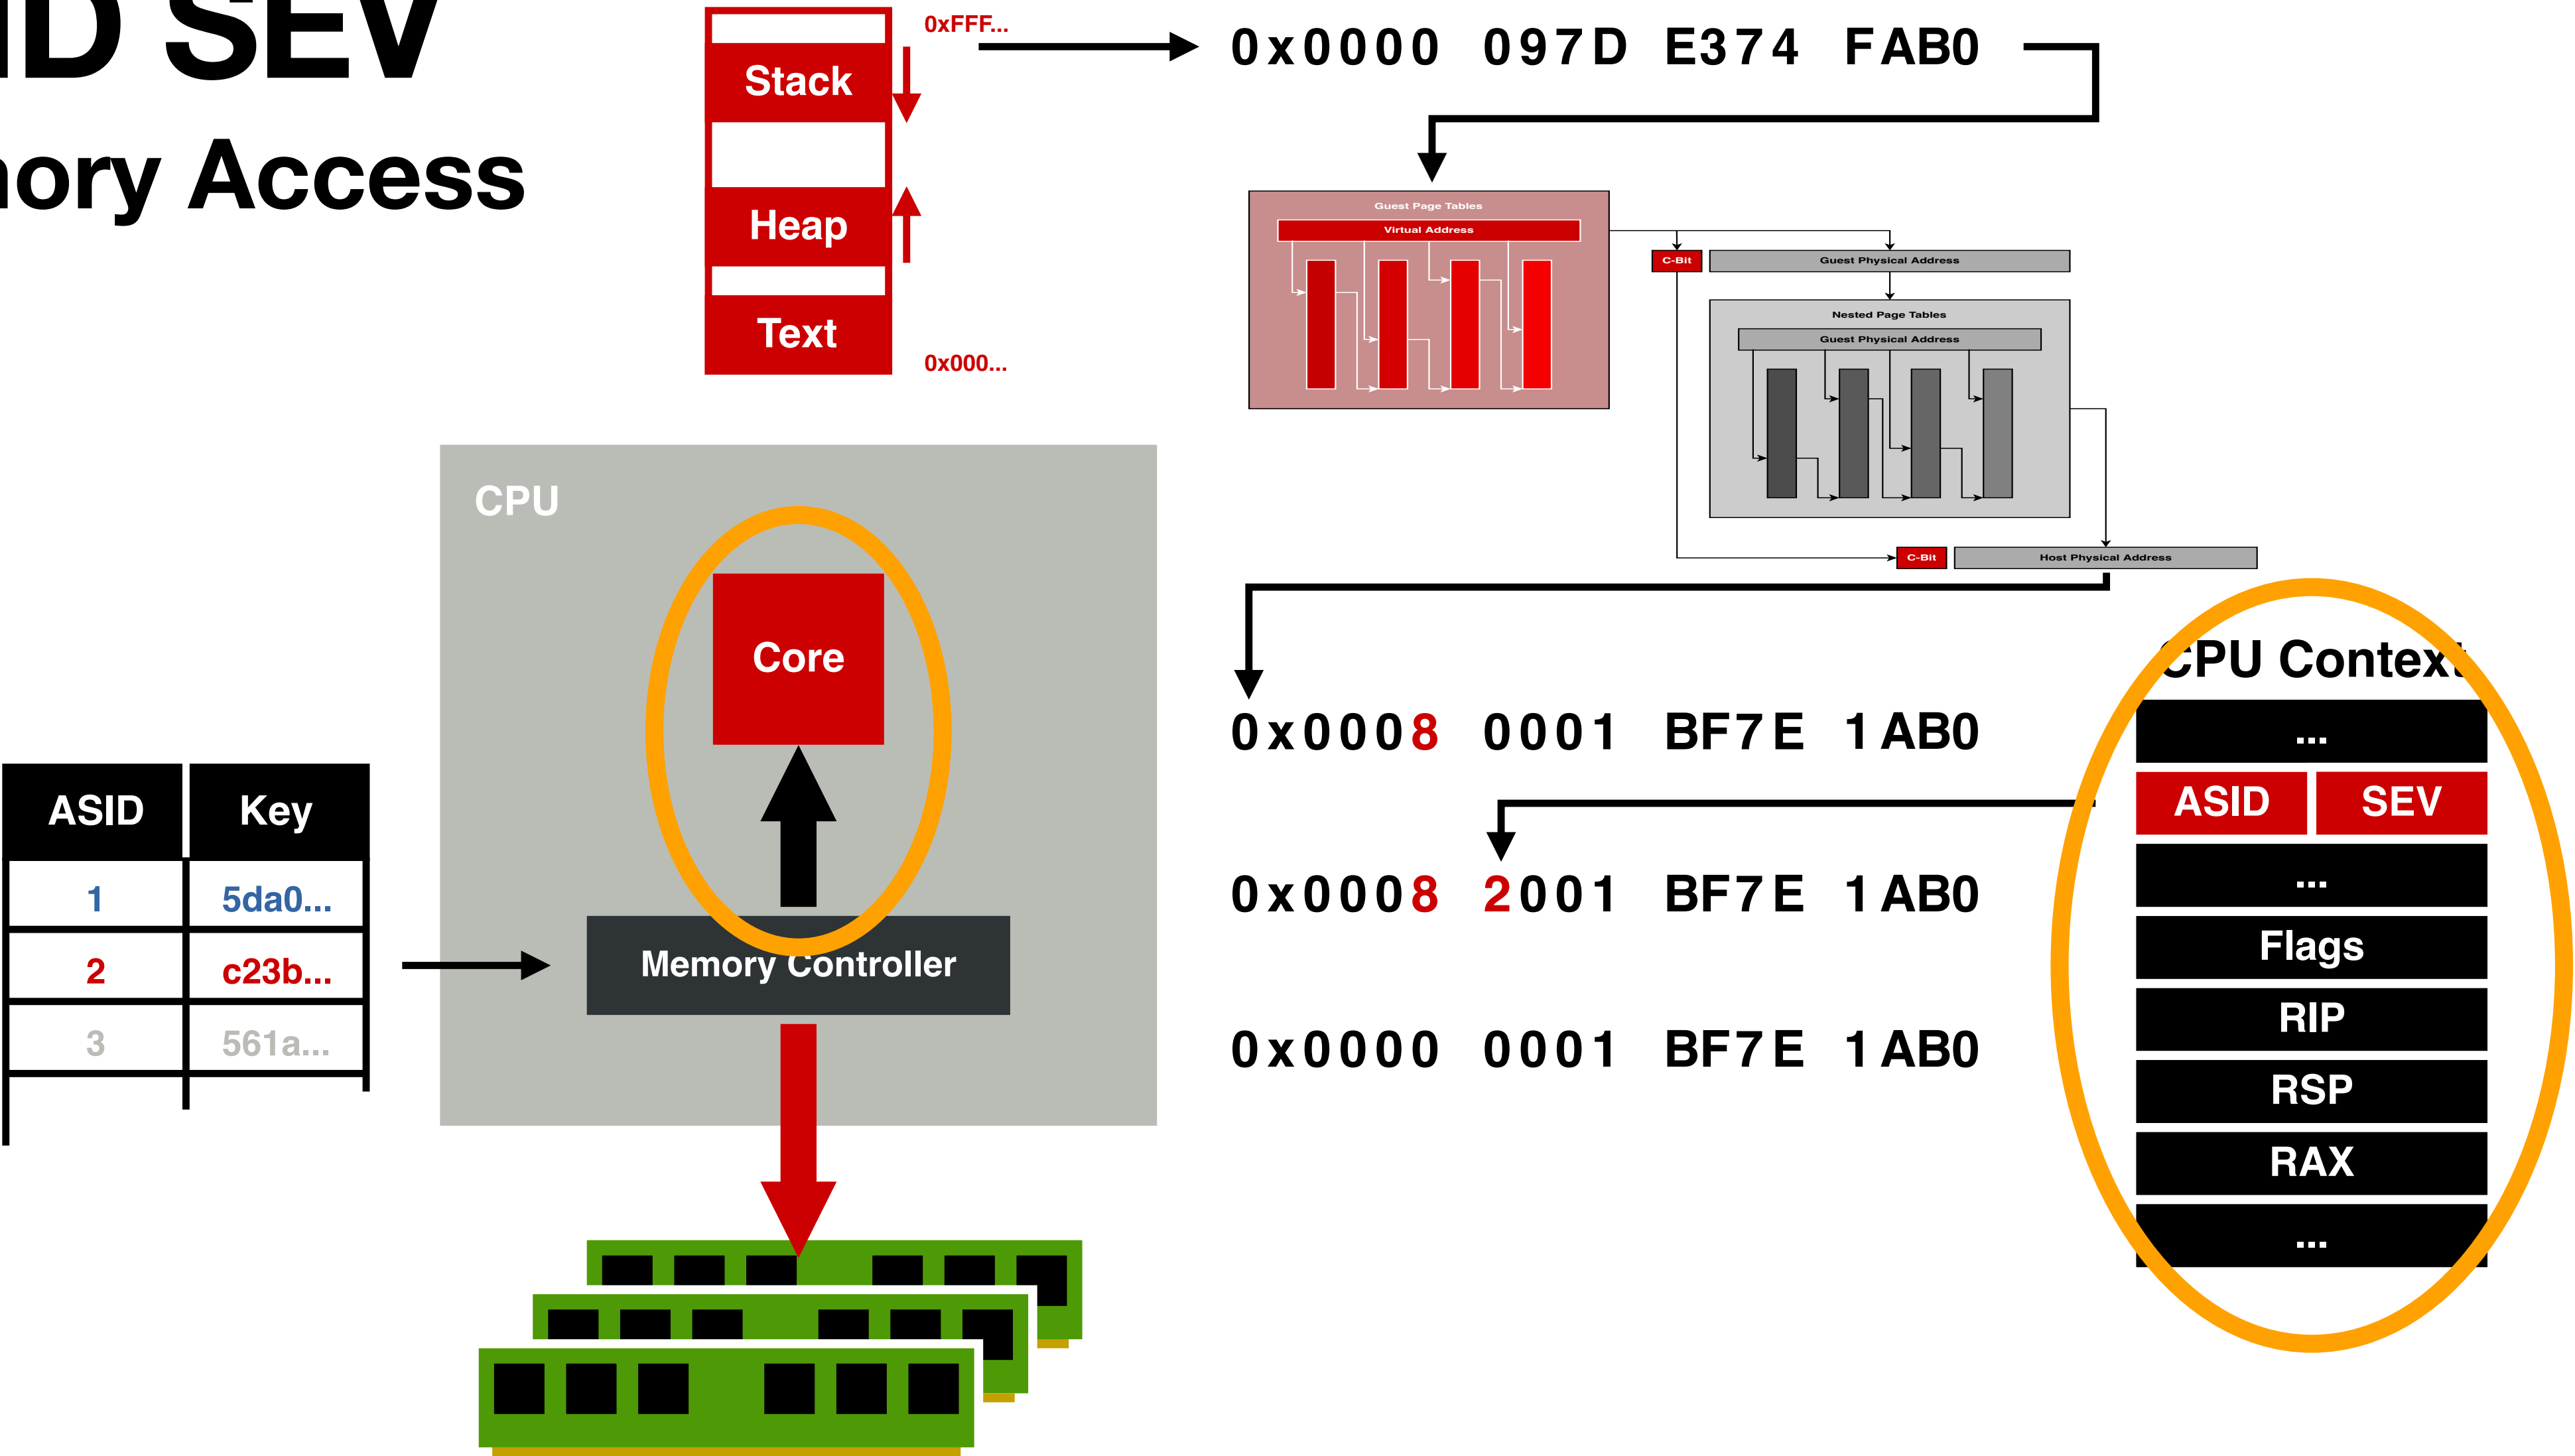
AMD SEV

Memory Access



AMD SEV

Memory Access



AMD SEV

Limitations

- Problem:
 - vCPU state visible to (untrusted) hypervisor
 - Including extended FPU state (AES-NI)
- Solution:
 - SEV-ES
 - Encrypting vCPU state

Confidential Computing with OpenBSD

Agenda

- Introduction
- First step: Memory encryption for VMs — SEV
- **Next step: vCPU state encryption — SEV-ES**
- Conclusion

AMD SEV-ES

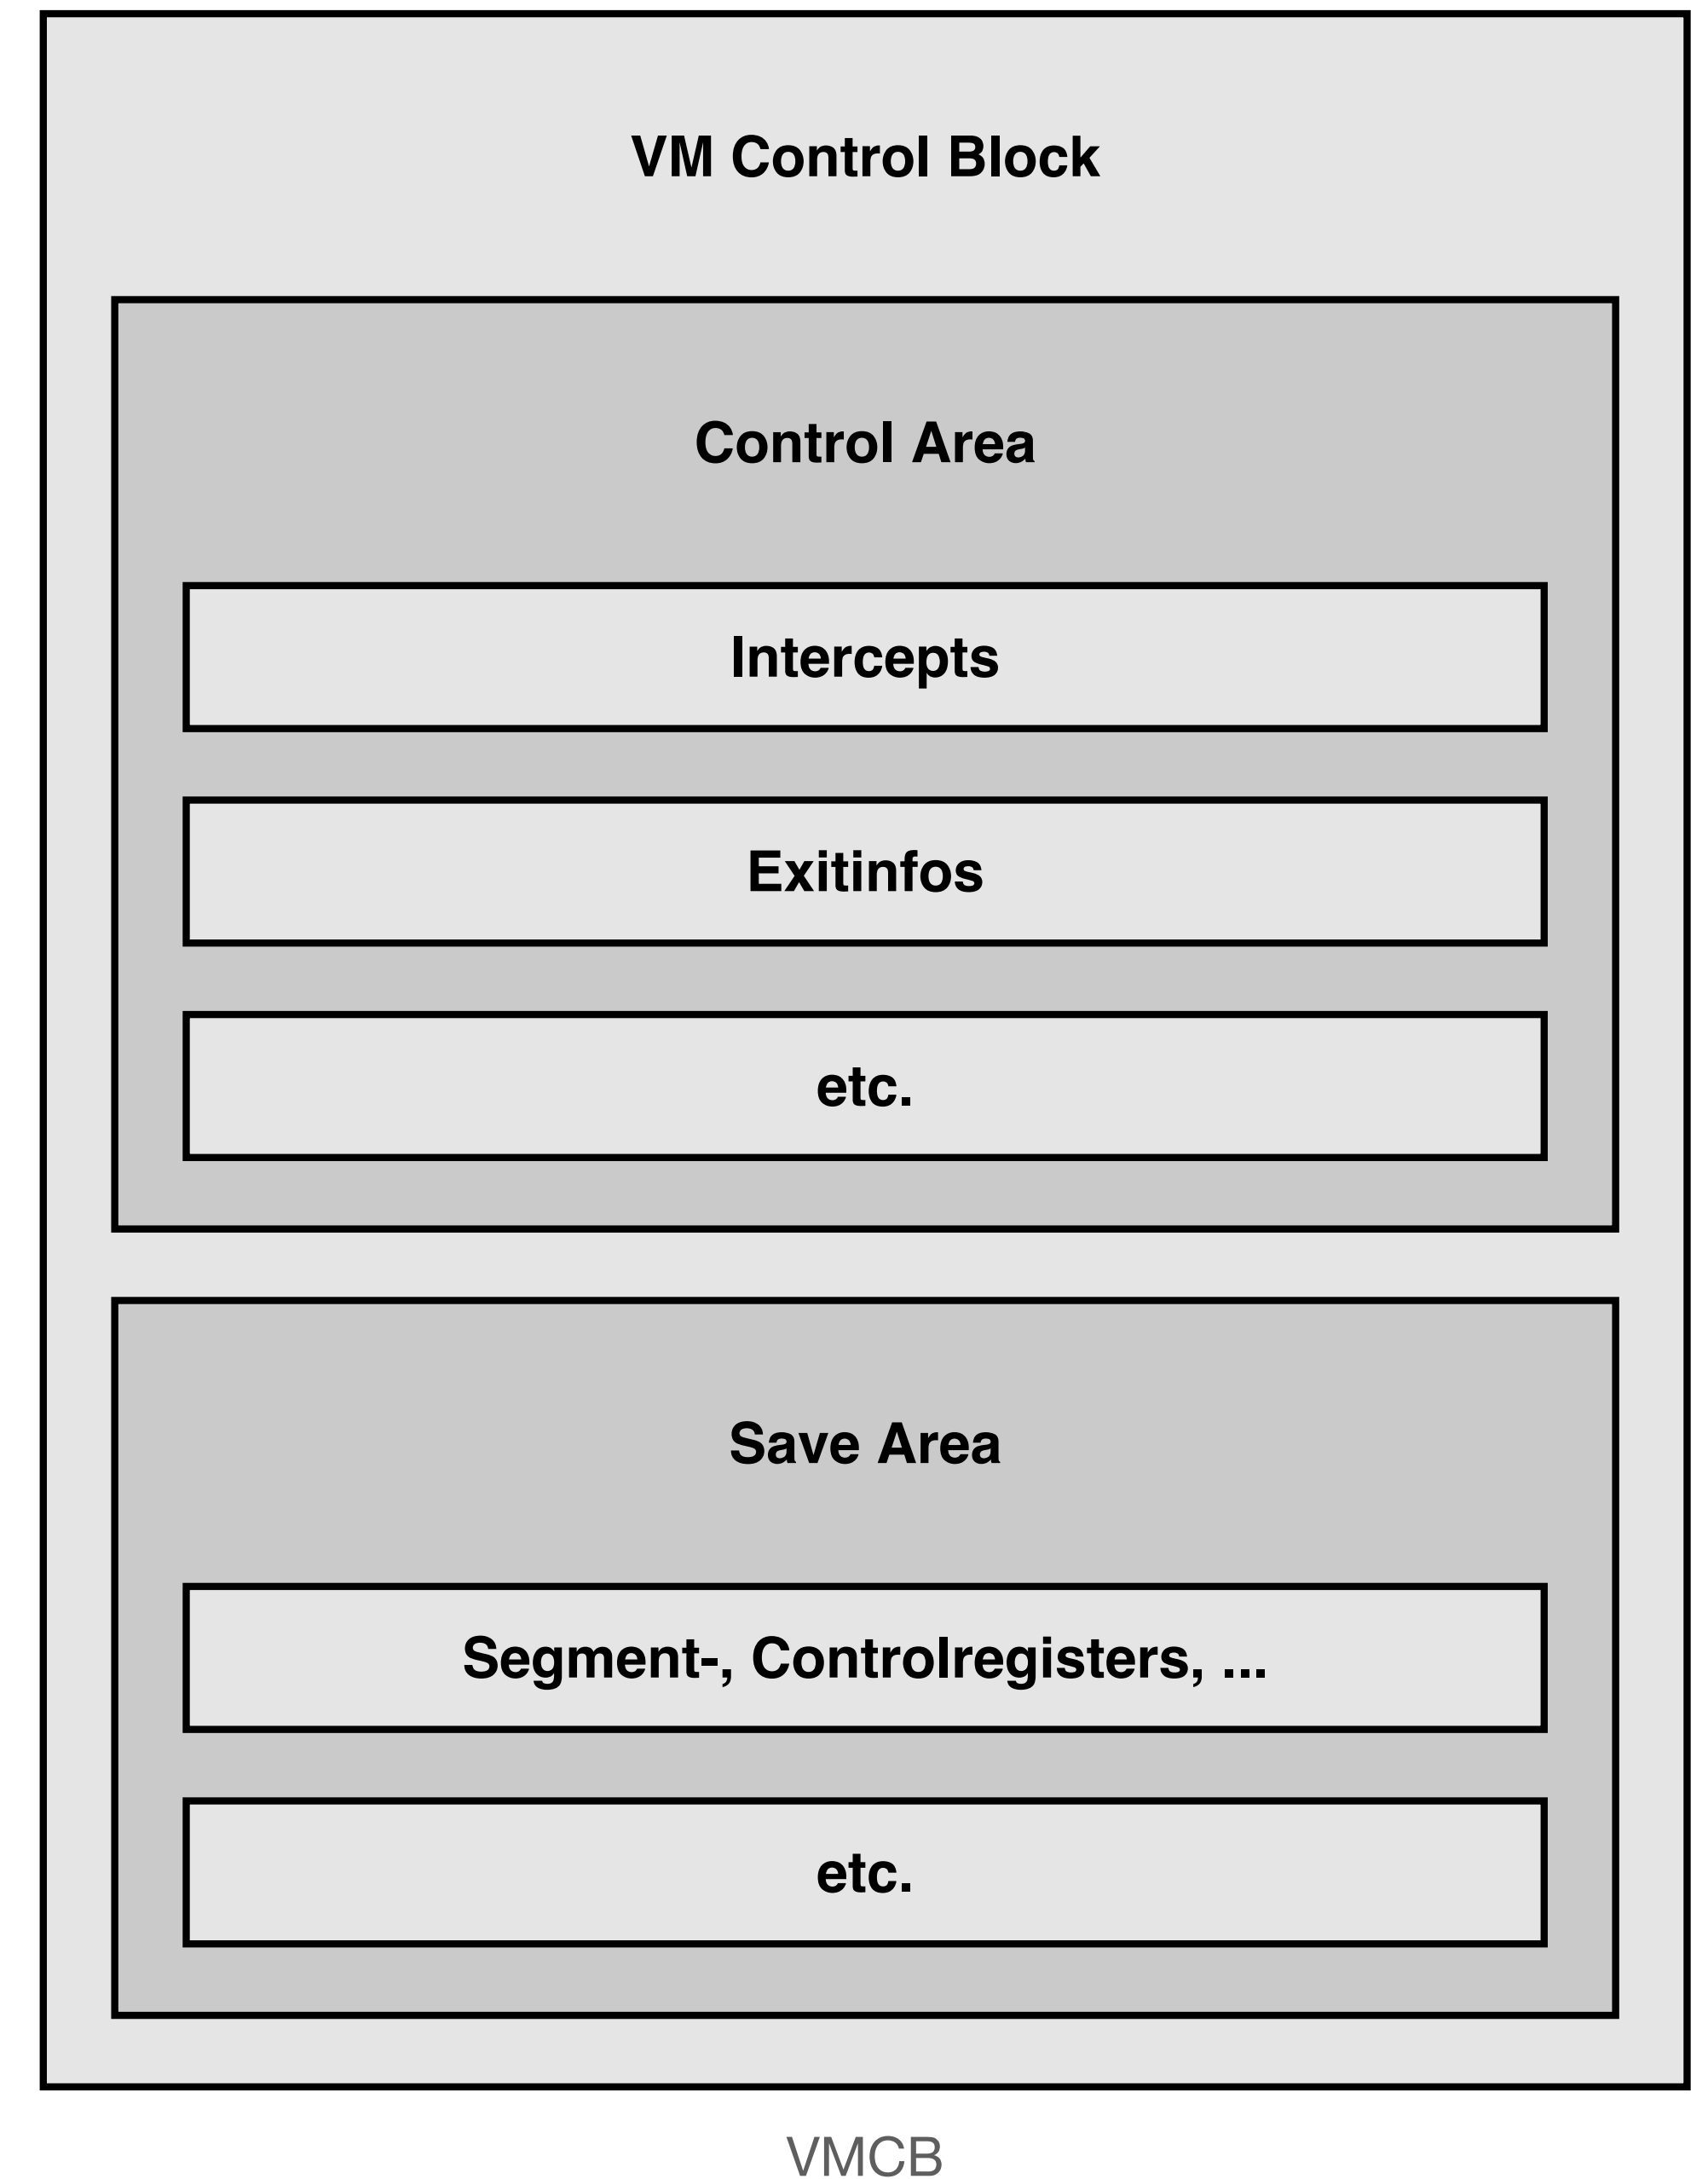
Encrypted vCPU State

- Regular SVM or SEV enabled VM:
 - Minimal state saved in VMCB
 - vmm(4) saves all remaining state in vCPU data structure
 - See exception/interrupt handling and stack frame
- SEV-ES enabled VM:
 - Full vCPU state saved automatically to encrypted VMSEA
 - vCPU state invisible (encrypted) for vmm(4)
- Host state saved to Host Save Area

AMD SEV-ES

VMCB

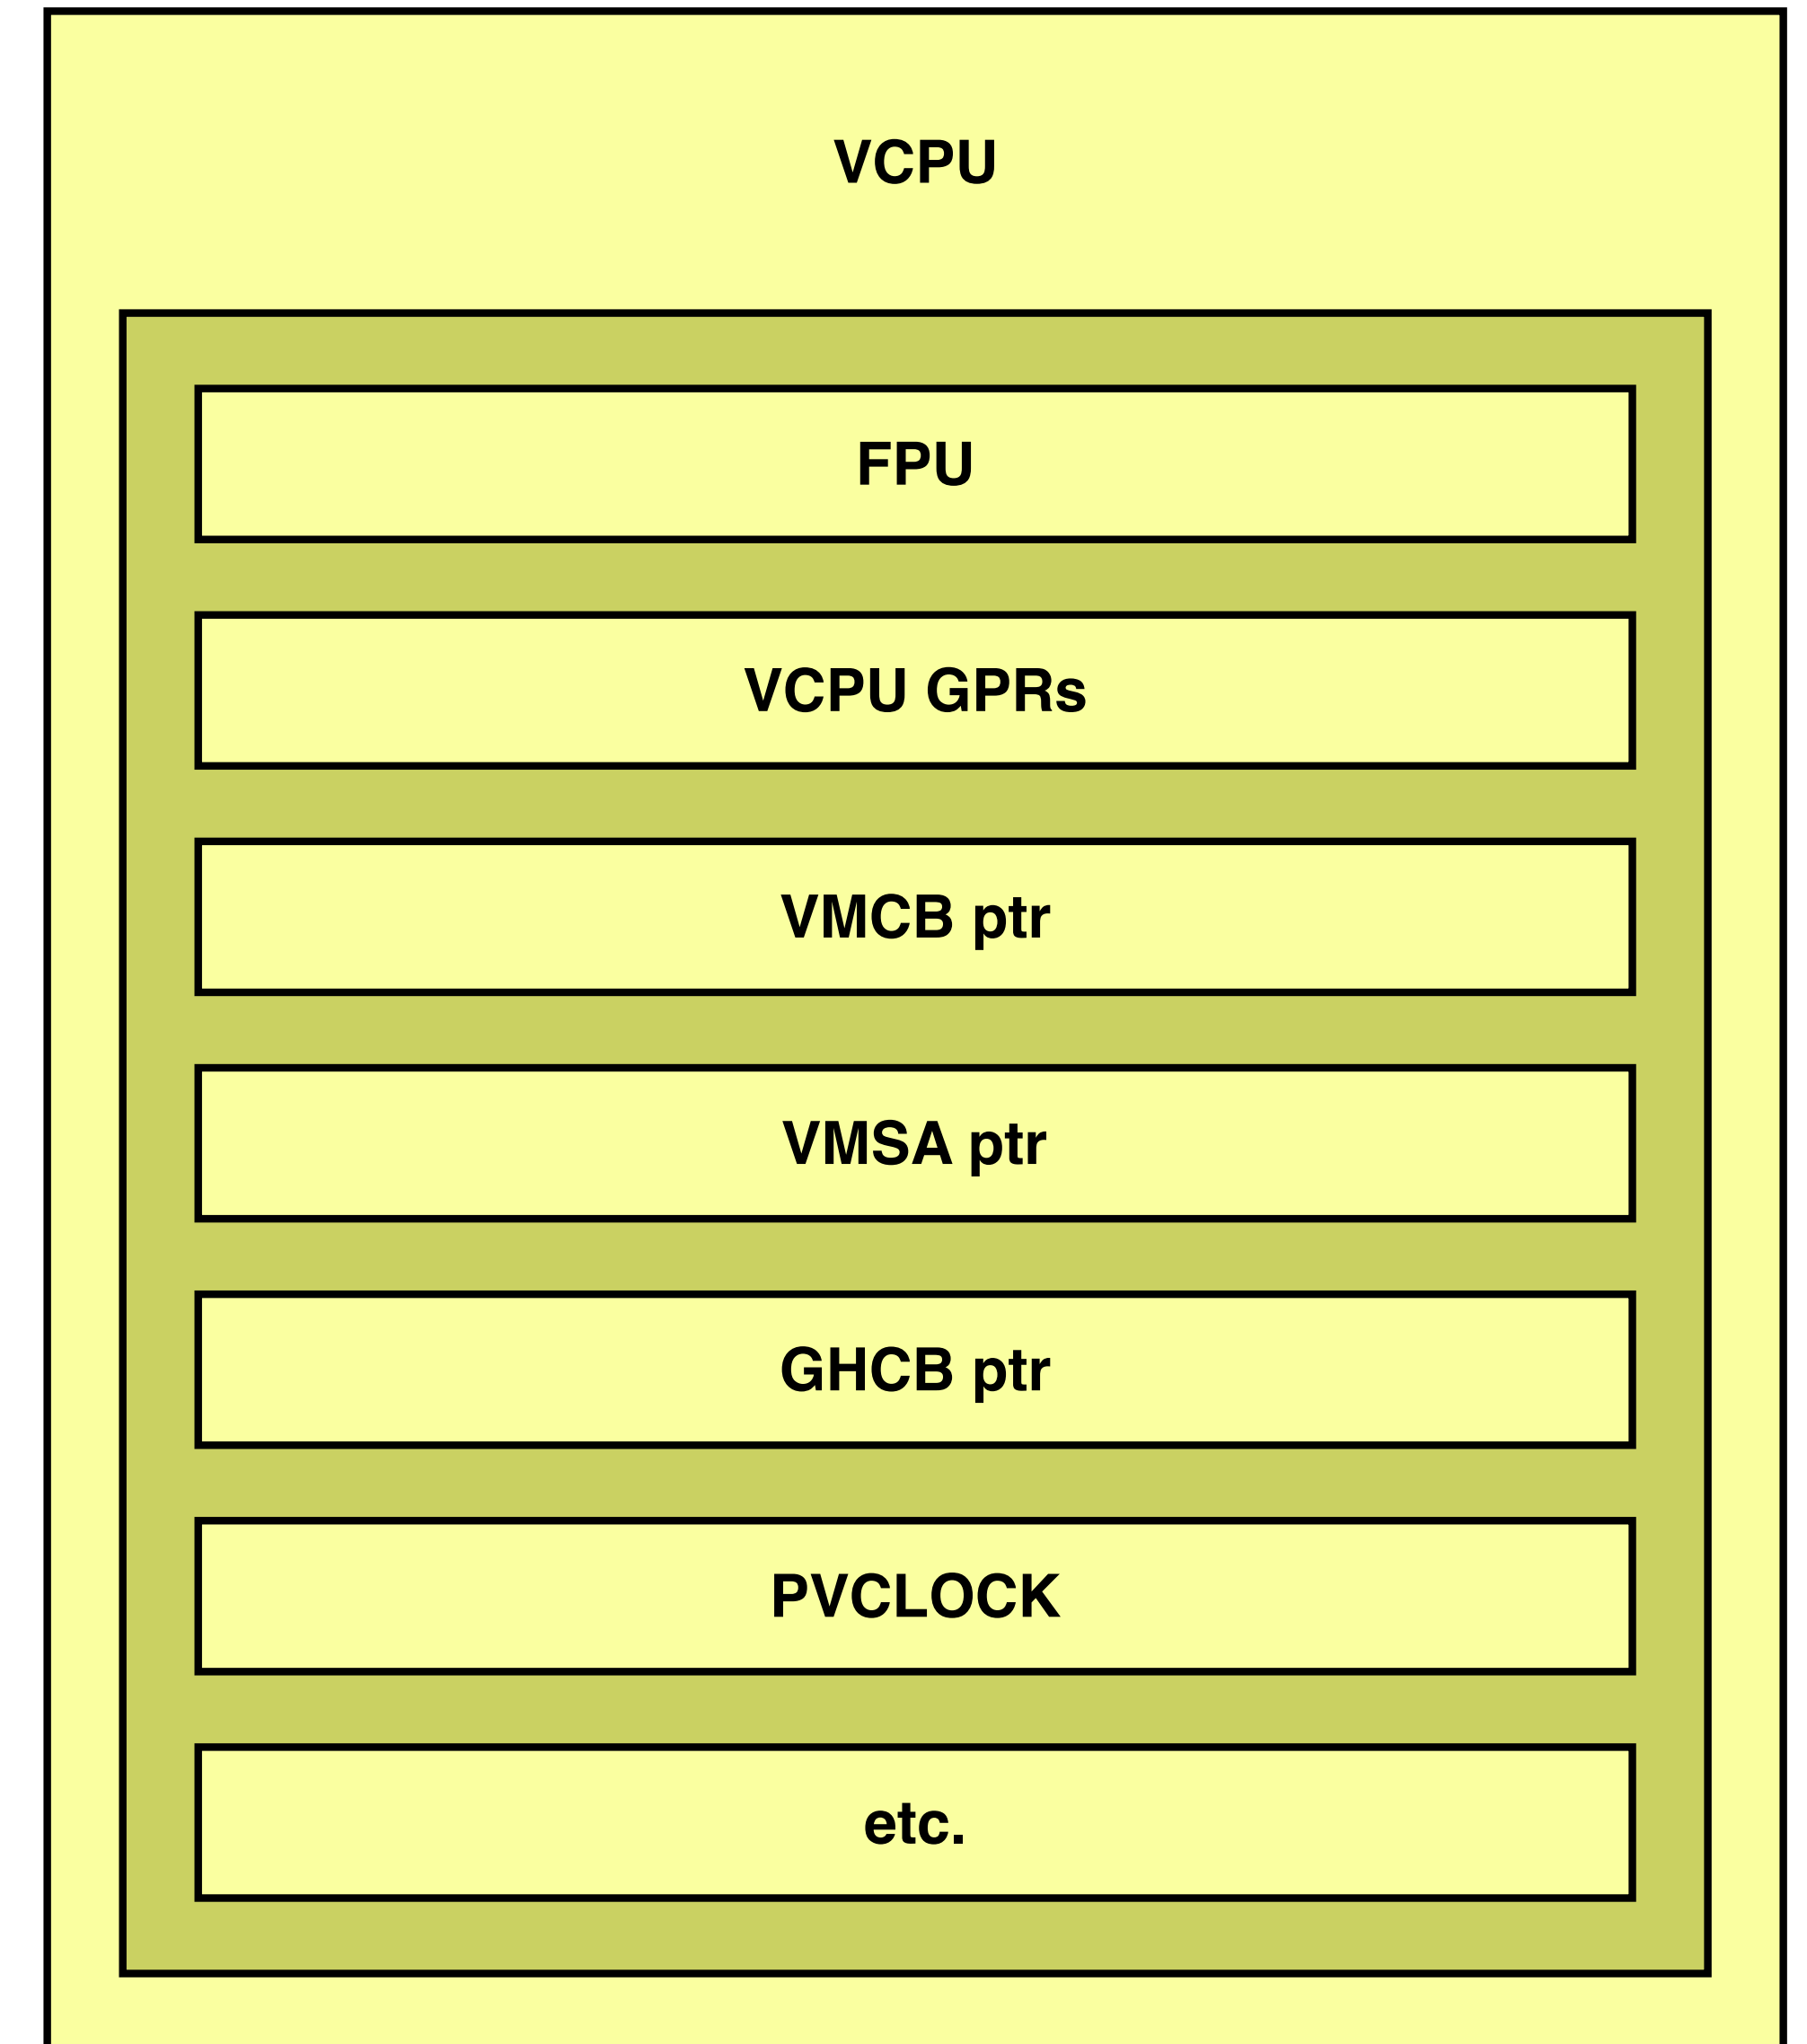
- Virtual Machine Control Block (VMCB)
 - Control Area
 - Minimal vCPU state
 - VMSAVE and VMLOAD



AMD SEV-ES

vCPU

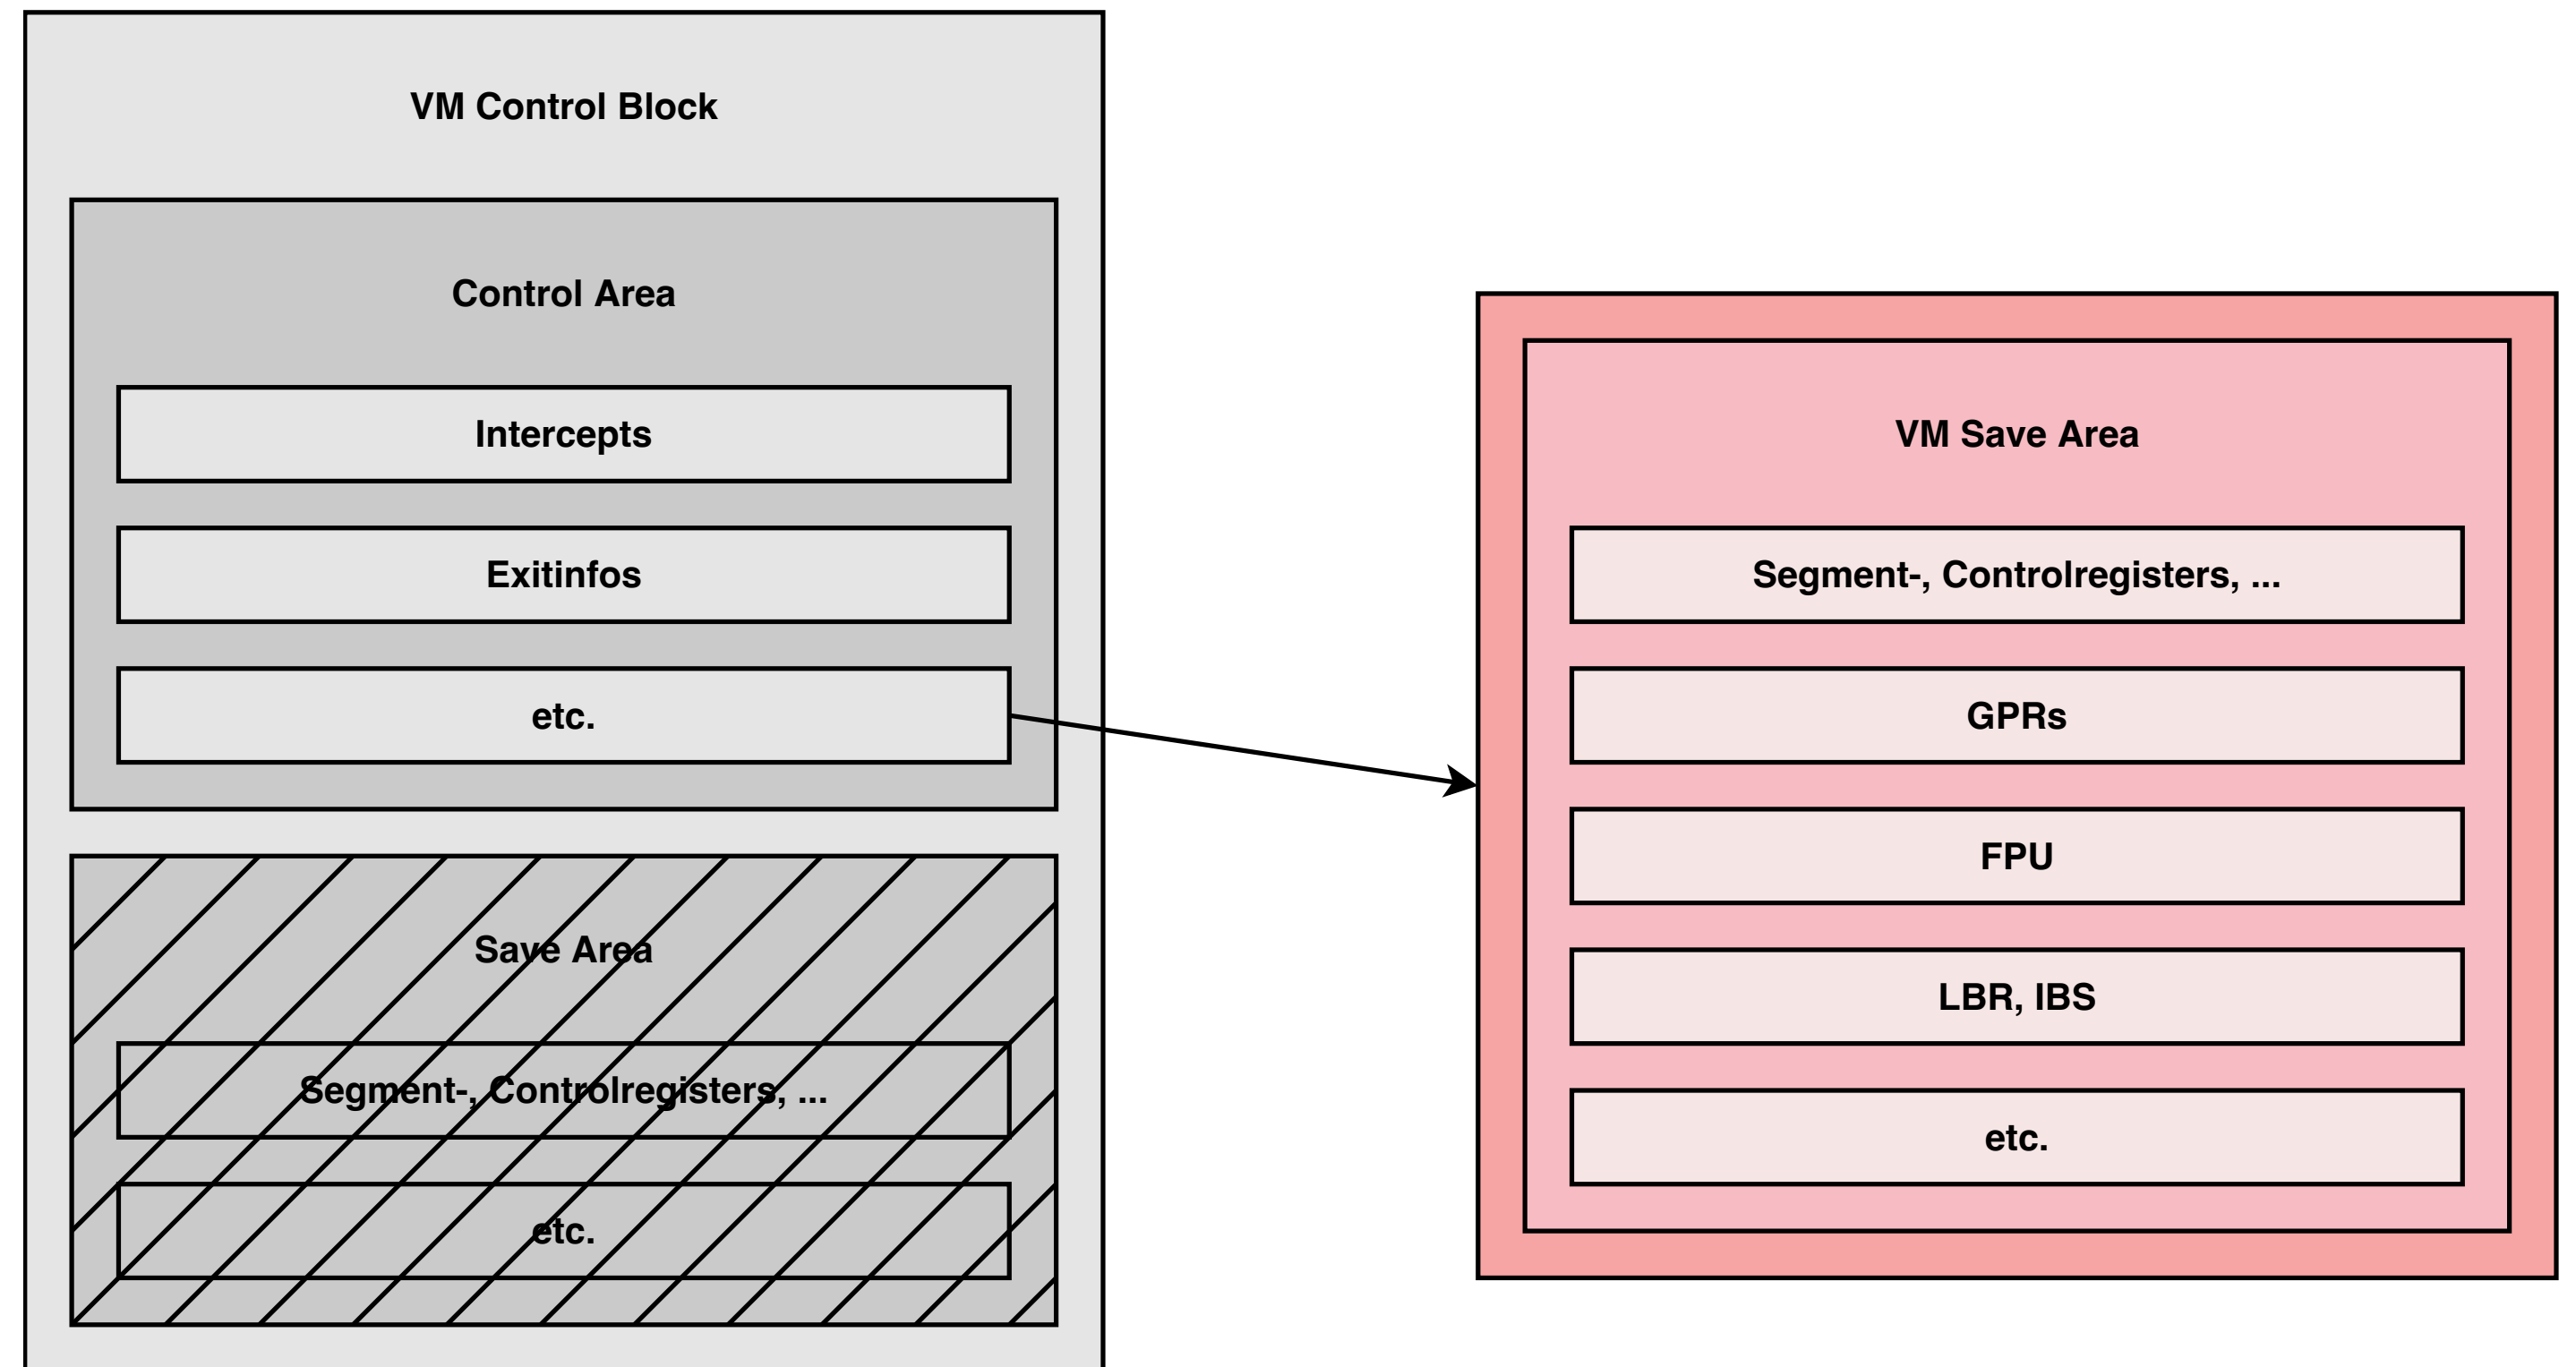
- vCPU data structure
 - Maintained by vmm(4)
- vCPU state
- Auxiliary data



AMD SEV-ES

VMMSA

- Virtual Machine Save Area
 - Maintained by CPU
 - Full vCPU state
 - Encrypted
- “Swapped” with host state

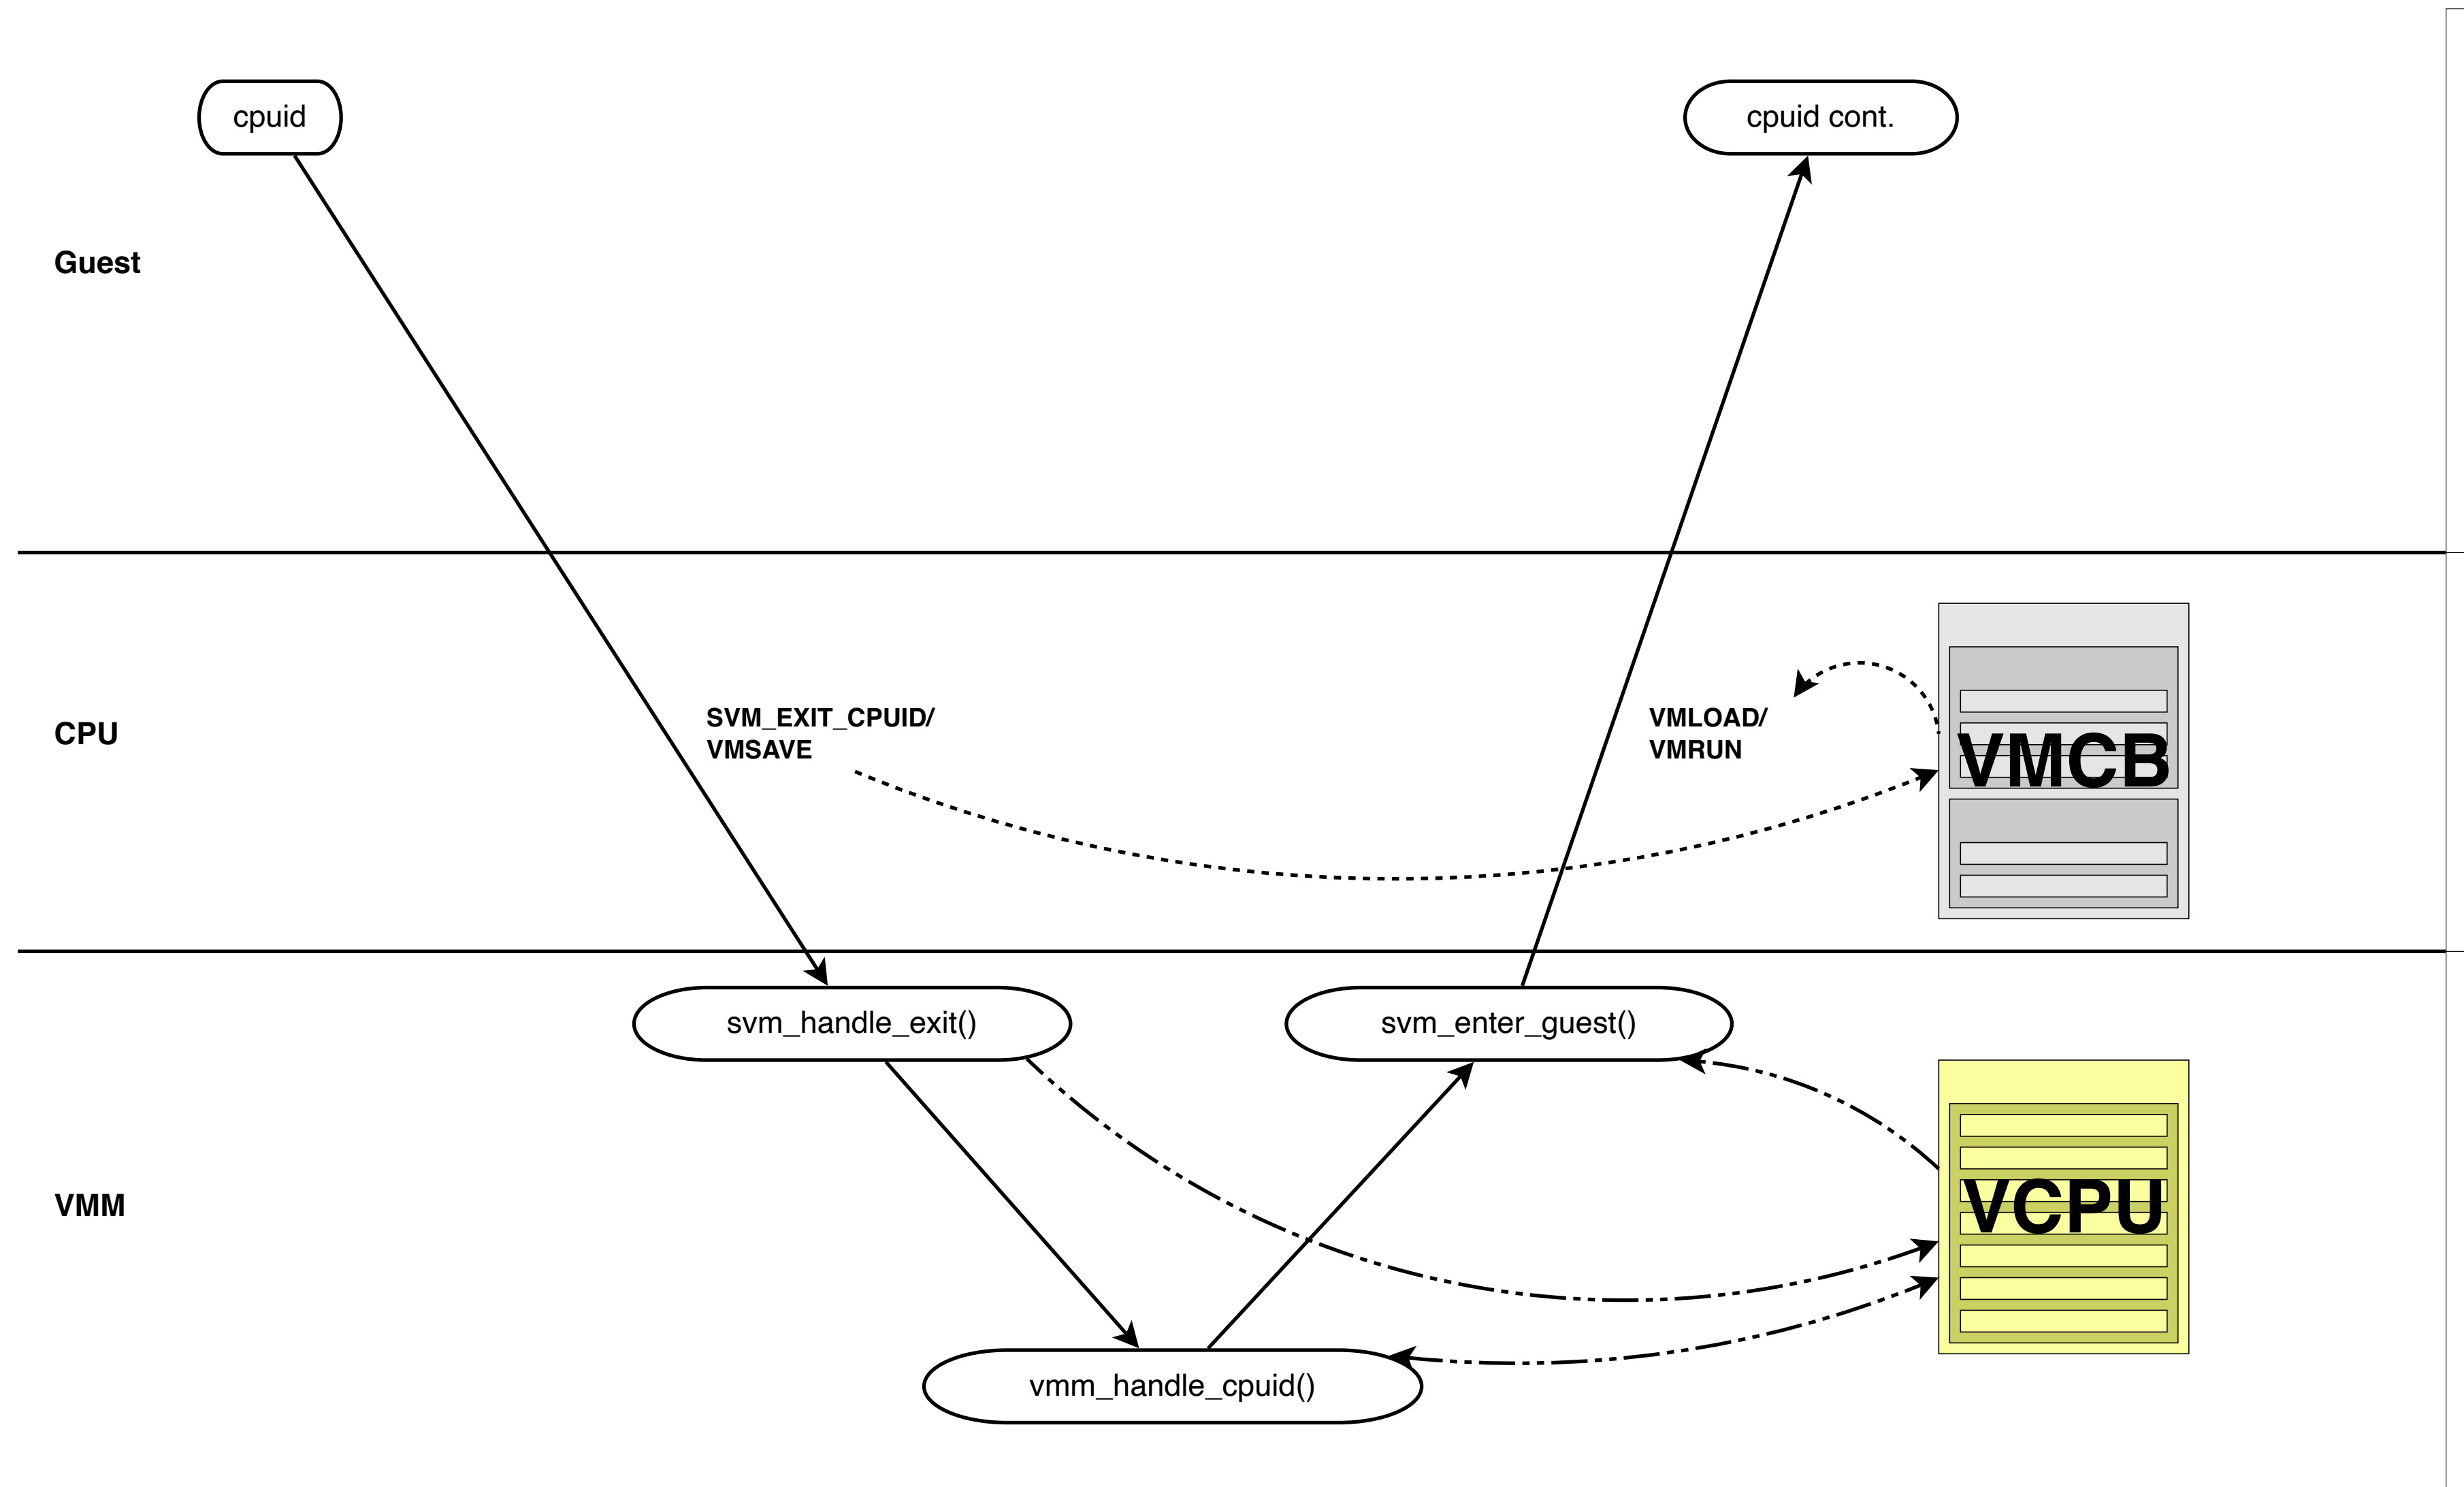


SVM and SEV

VM Exit and Entry

- VM Exit
 - Minimal state and hidden state saved to VMCB with VMSAVE
 - vCPU state saved by vmm(4)
- VM Entry
 - vCPU state restored by vmm(4)
 - State in VMCB restored with VMLOAD

VM Exit

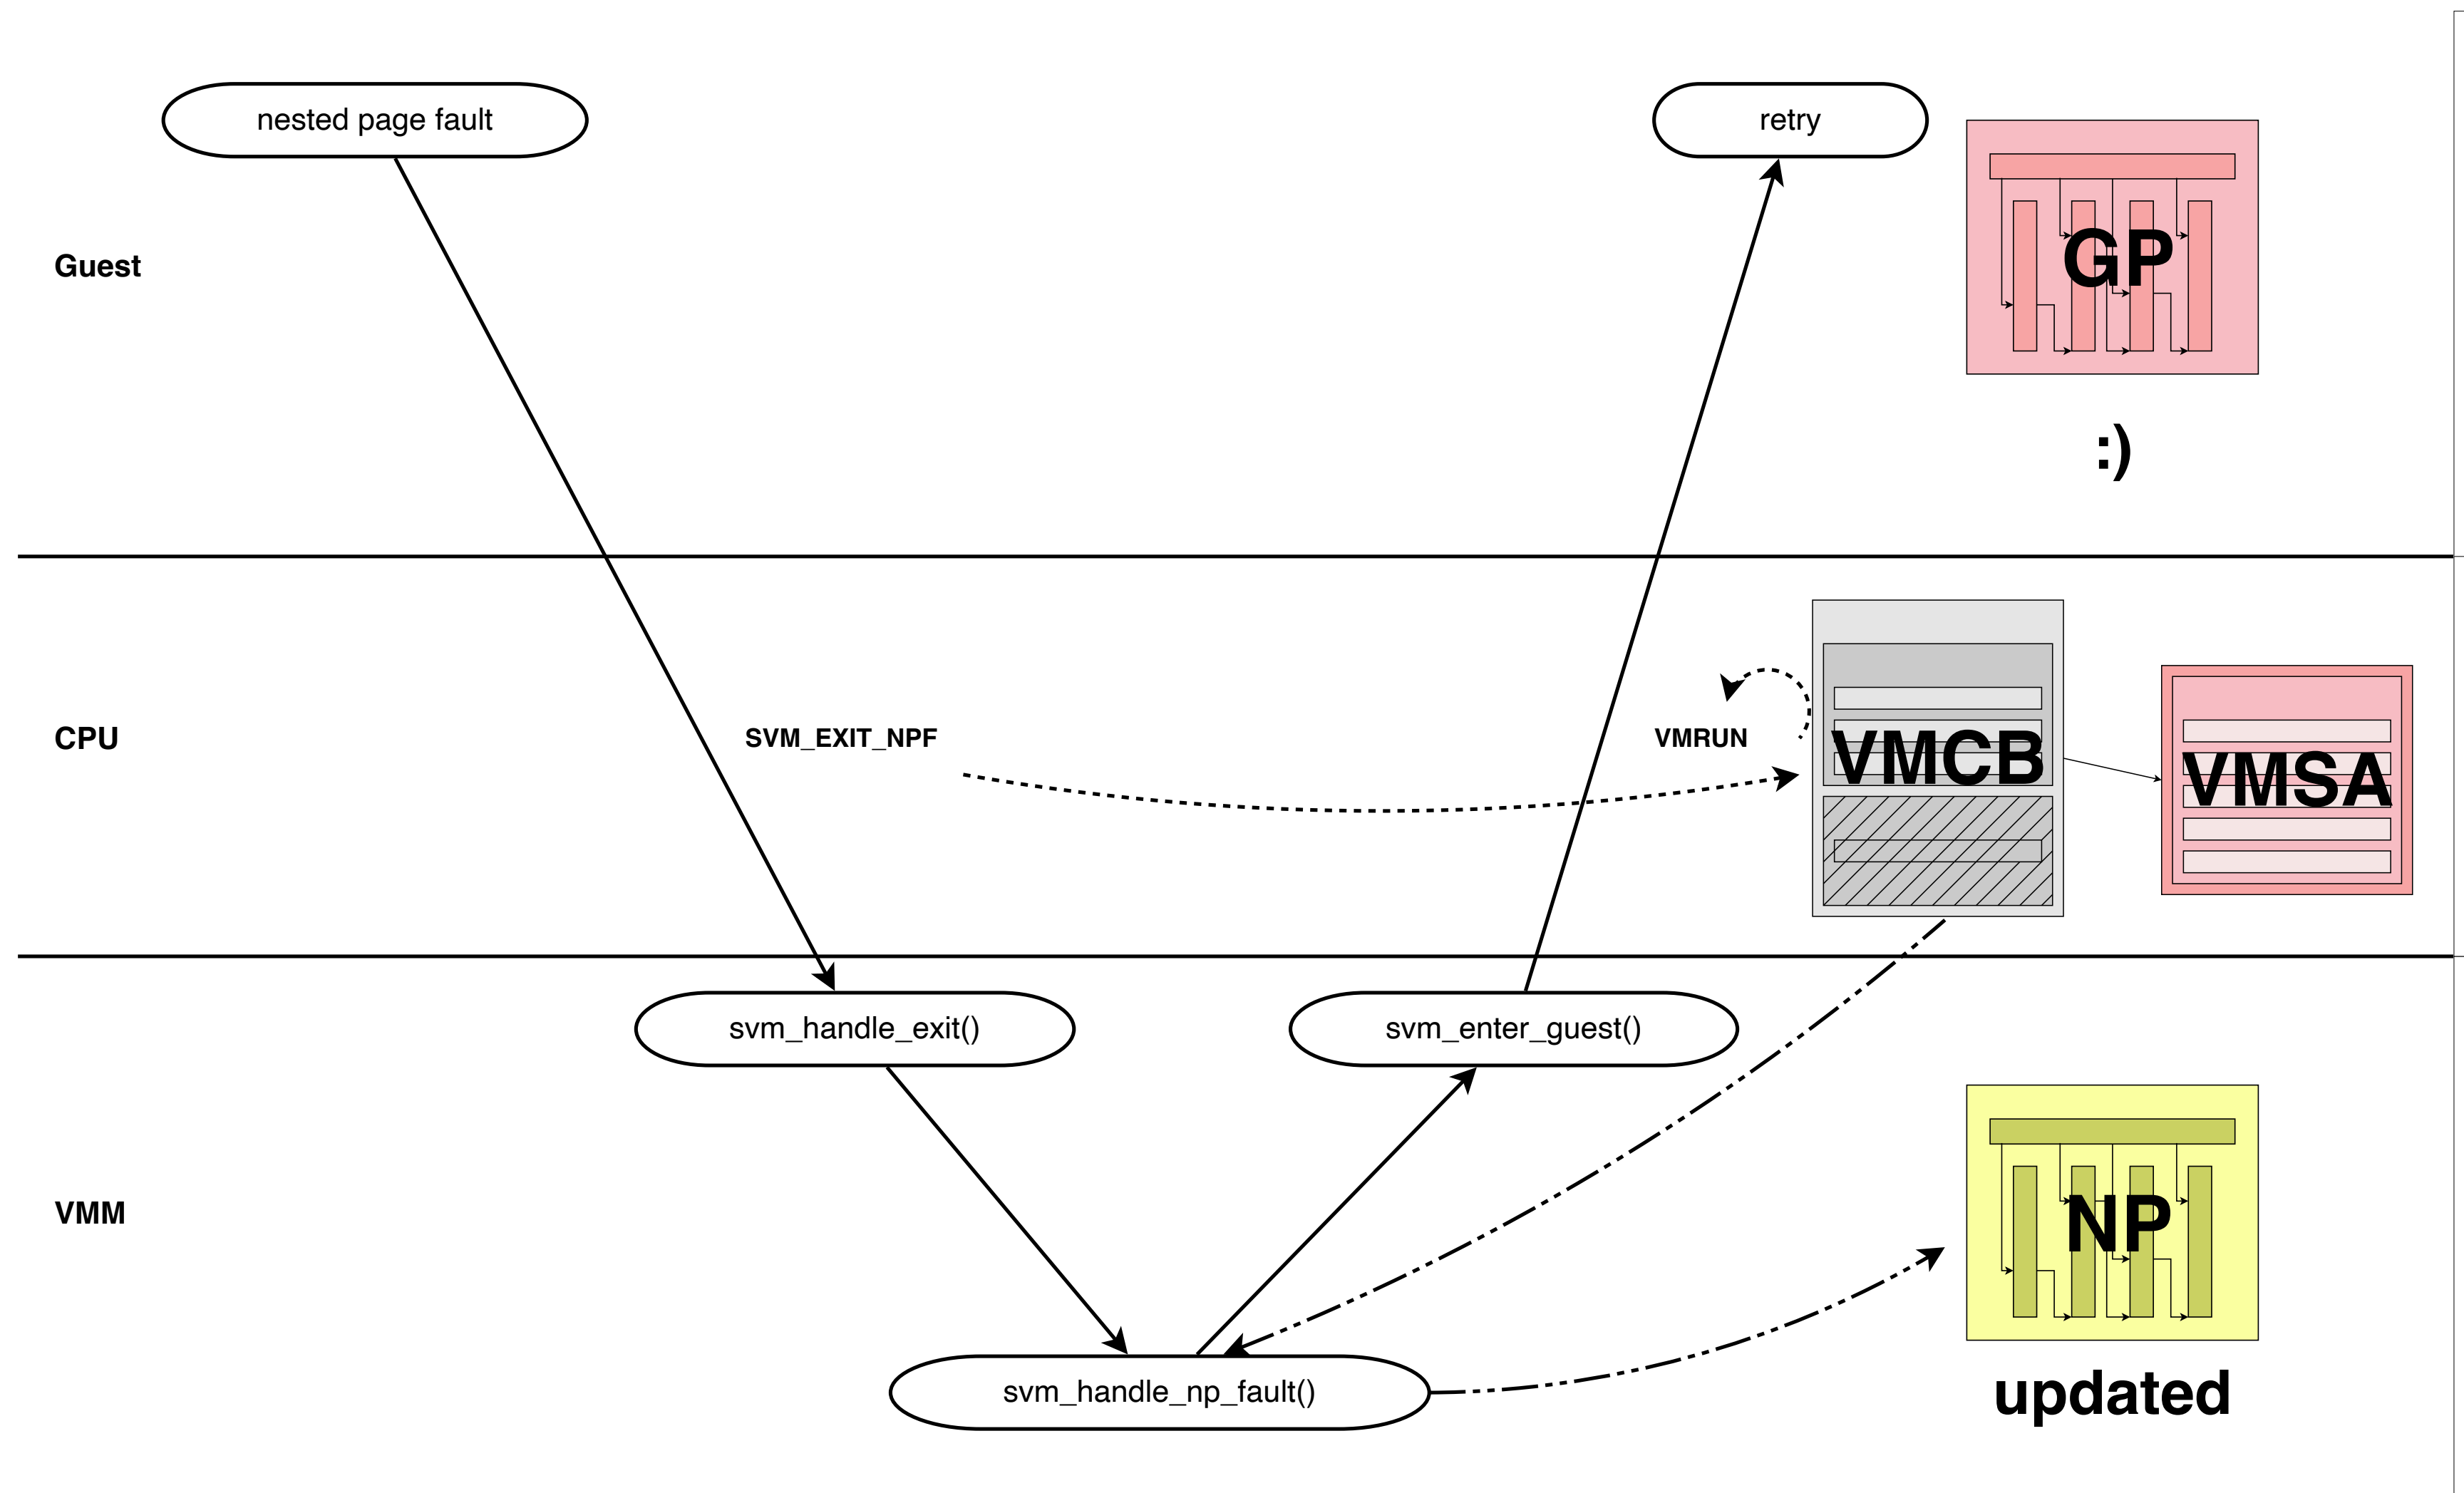


AMD SEV-ES

VM Exit Types

- Automatic Exits
 - Asynchronous
 - No vCPU state needed by vmm(4)
- Non-Automatic Exits
 - All other exits
 - Guest decides on what vCPU state to expose

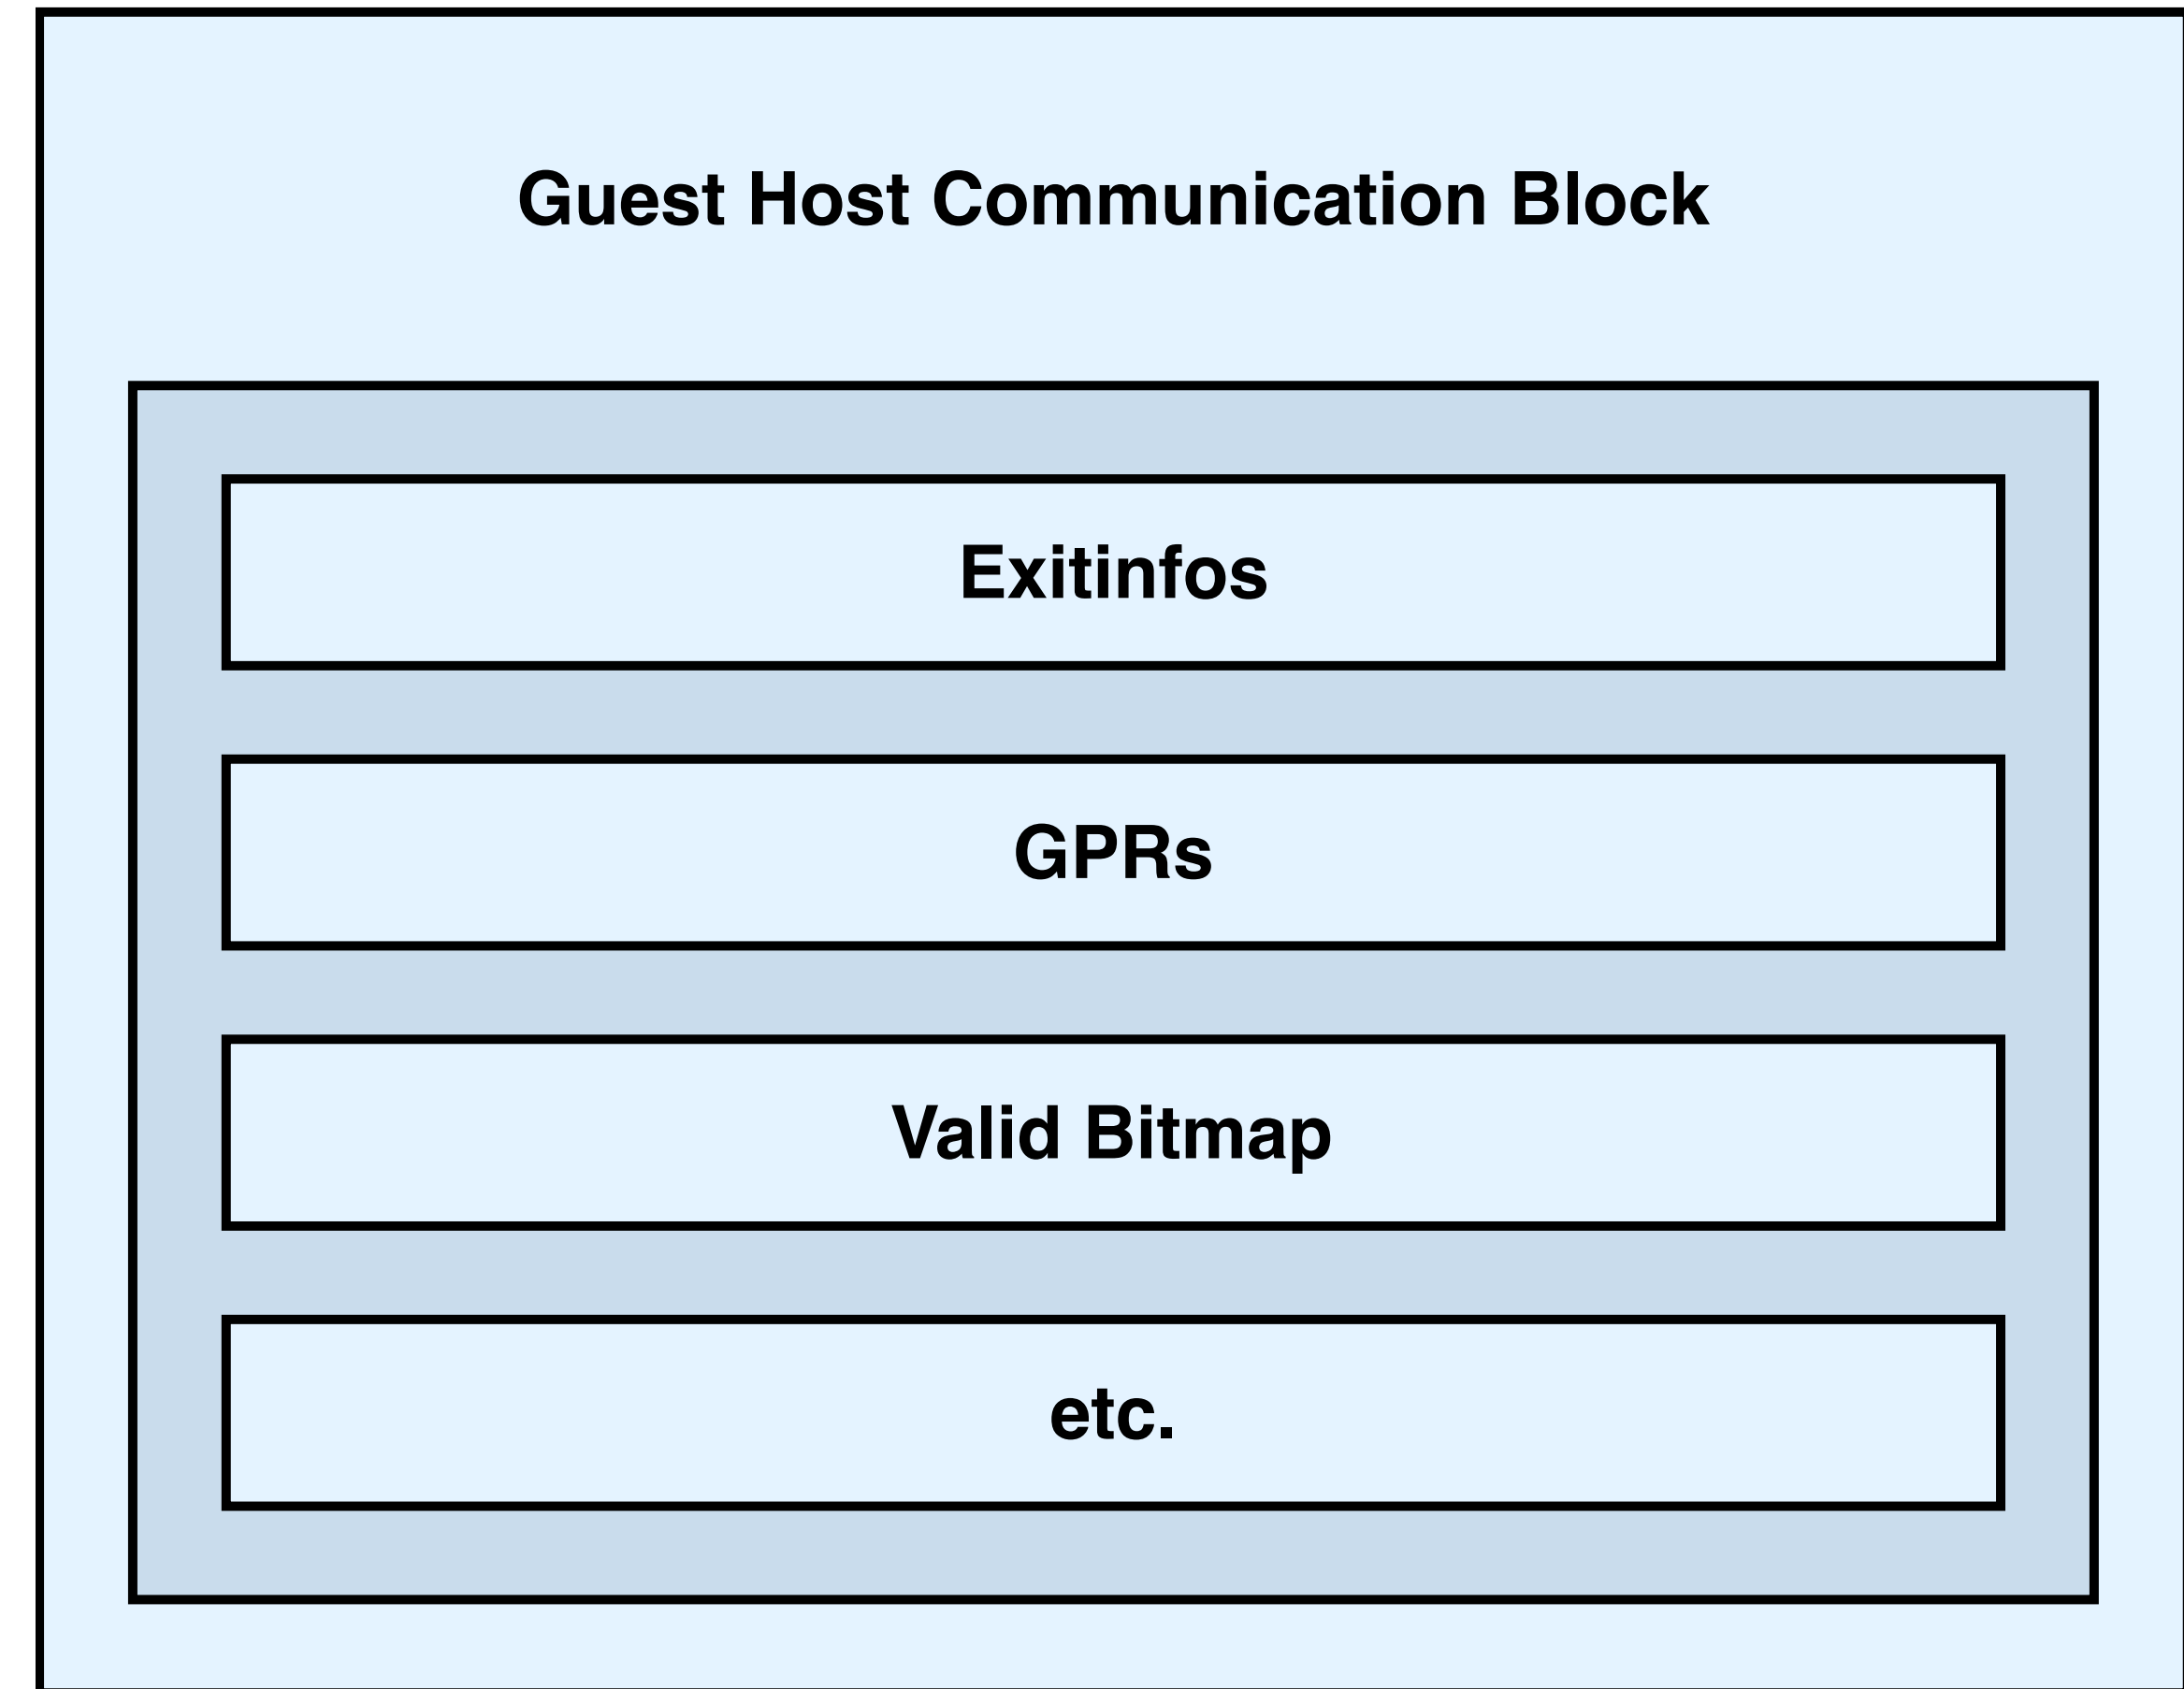
AE VM Exit



#VC Trap

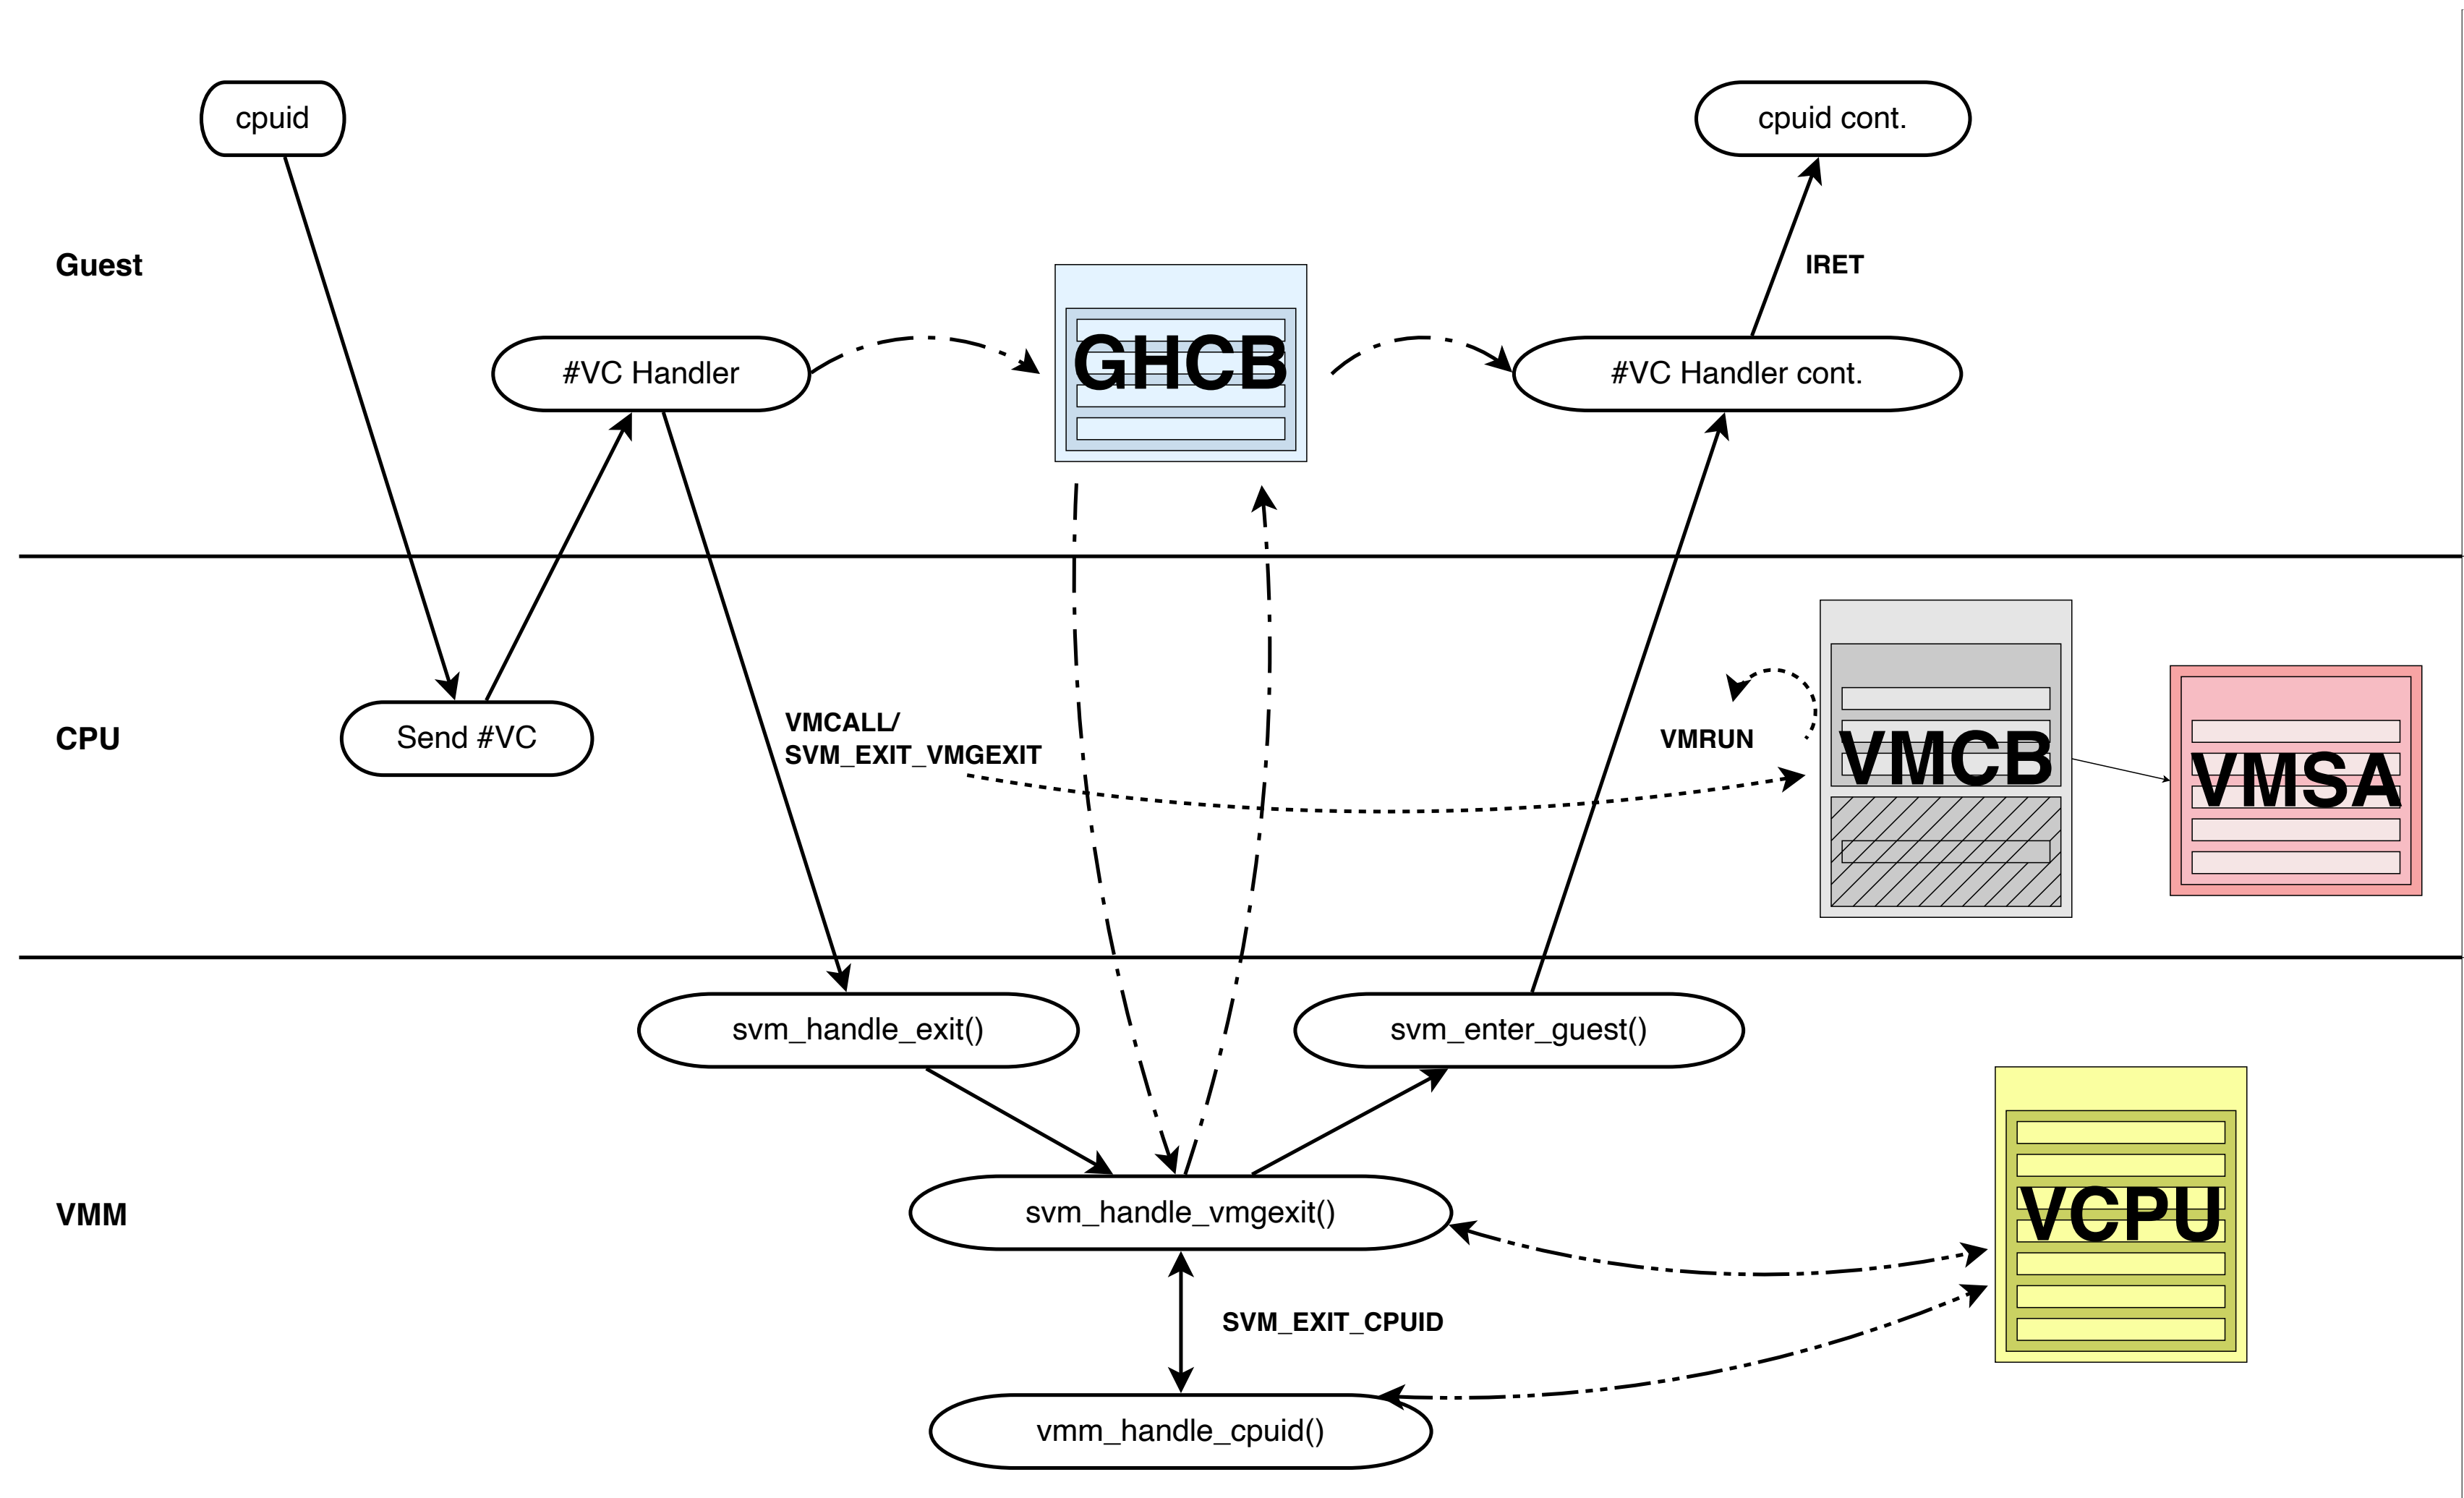
Non-Automatic VM Exits

- VM Exit redirected to #VC trap handler
- Guest decides on what vCPU to be shared with vmm(4)
- Software defined Guest Host Communication Block (GHCB)
- Unencrypted memory shared by guest with vmm(4)
- GHCB MSR points to GPA of GHCB



GHCB

NAE VM Exit



SEV-ES Bootstrap

*

#VC in locore0

- Challenge:
 - CPUID might raise #VC
 - Guest is not “enlightened” yet
 - Plain GENERIC kernel
- Tentative #VC handler:
 - No SEV-ES, nothing happens, all fine :)
 - SEV-ES enabled guest:
 - Handle #VC trap

SEV-ES Bootstrap

*

#VC in locore0

- #VC handler:
 - Paging not enabled yet
 - No GHCB shared with vmm(4)
- GHCB MSR protocol:
 - Use low 12 bits to encode requests
 - WRMSR followed by VMCALL/VMGEXIT
 - vmm(4) encodes response as GHCB GPA in VMCB (@0xA0)
 - RDMSR by guest

Paravirtualisation

*

Avoiding #VC

- vmm(4) emulates PIC i8259
 - OUT in IRQ handler
 - Each IRQ raises several #VC
- Paravirtualising IN/OUT
 - When SEV-ES is enabled
 - Codepatch IN/OUT with paravirtualised version
 - Completely avoids #VC after bootstrapping the kernel



Confidential Computing with OpenBSD

Agenda

- Introduction
- First step: Memory encryption for VMs — SEV
- Next step: vCPU state encryption — SEV-ES
- **Conclusion**

Conclusion

Next step and beyond

- SEV-ES works with OpenBSD-current:
 - GENERIC supports both host and guest
 - psp(4), vmm(4) and vmd(8) support implemented/updated
 - Integration into source tree is work in progress
- Next goal:
 - Support for SEV-SNP
 - OpenBSD guest already running on Linux/KVM host
 - Performance?
 - Attestation?
 - Security?
 - Linux guest on OpenBSD?
 - ...

Thanks!
Questions?

Don't forget to remember!